

MF20006: Introduction to Computer Science

Lecture 2: Computer Architecture

Hui Xu

xuh@fudan.edu.cn



Outline

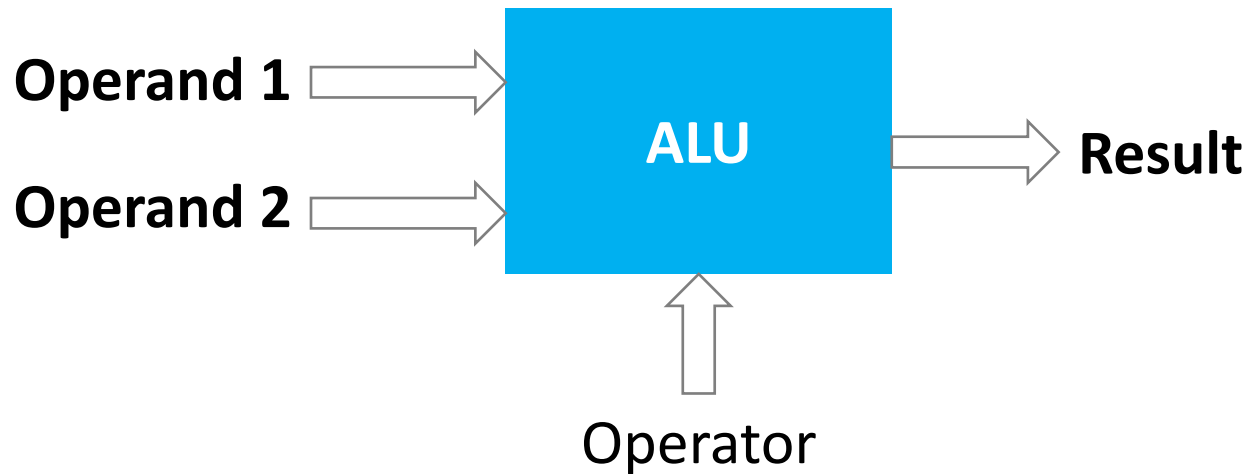
1. CPU and GPU

2. ISA and Assembly Code

3. I/O and More

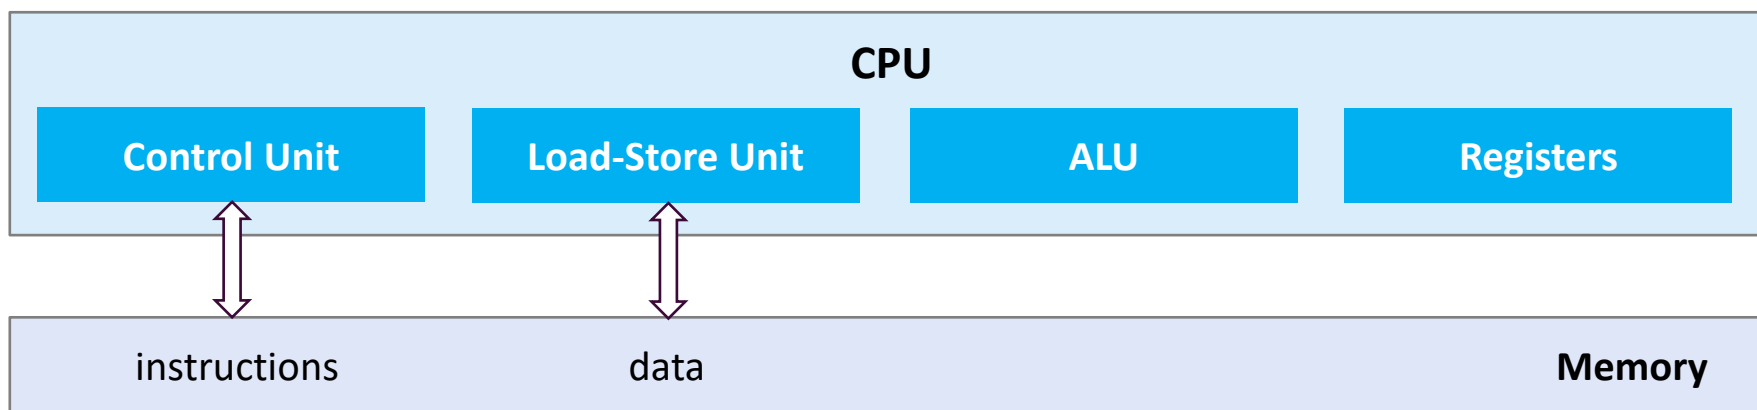
1. CPU and GPU

Recall: Transistor => Logic Gate => Arithmetic Logic Unit

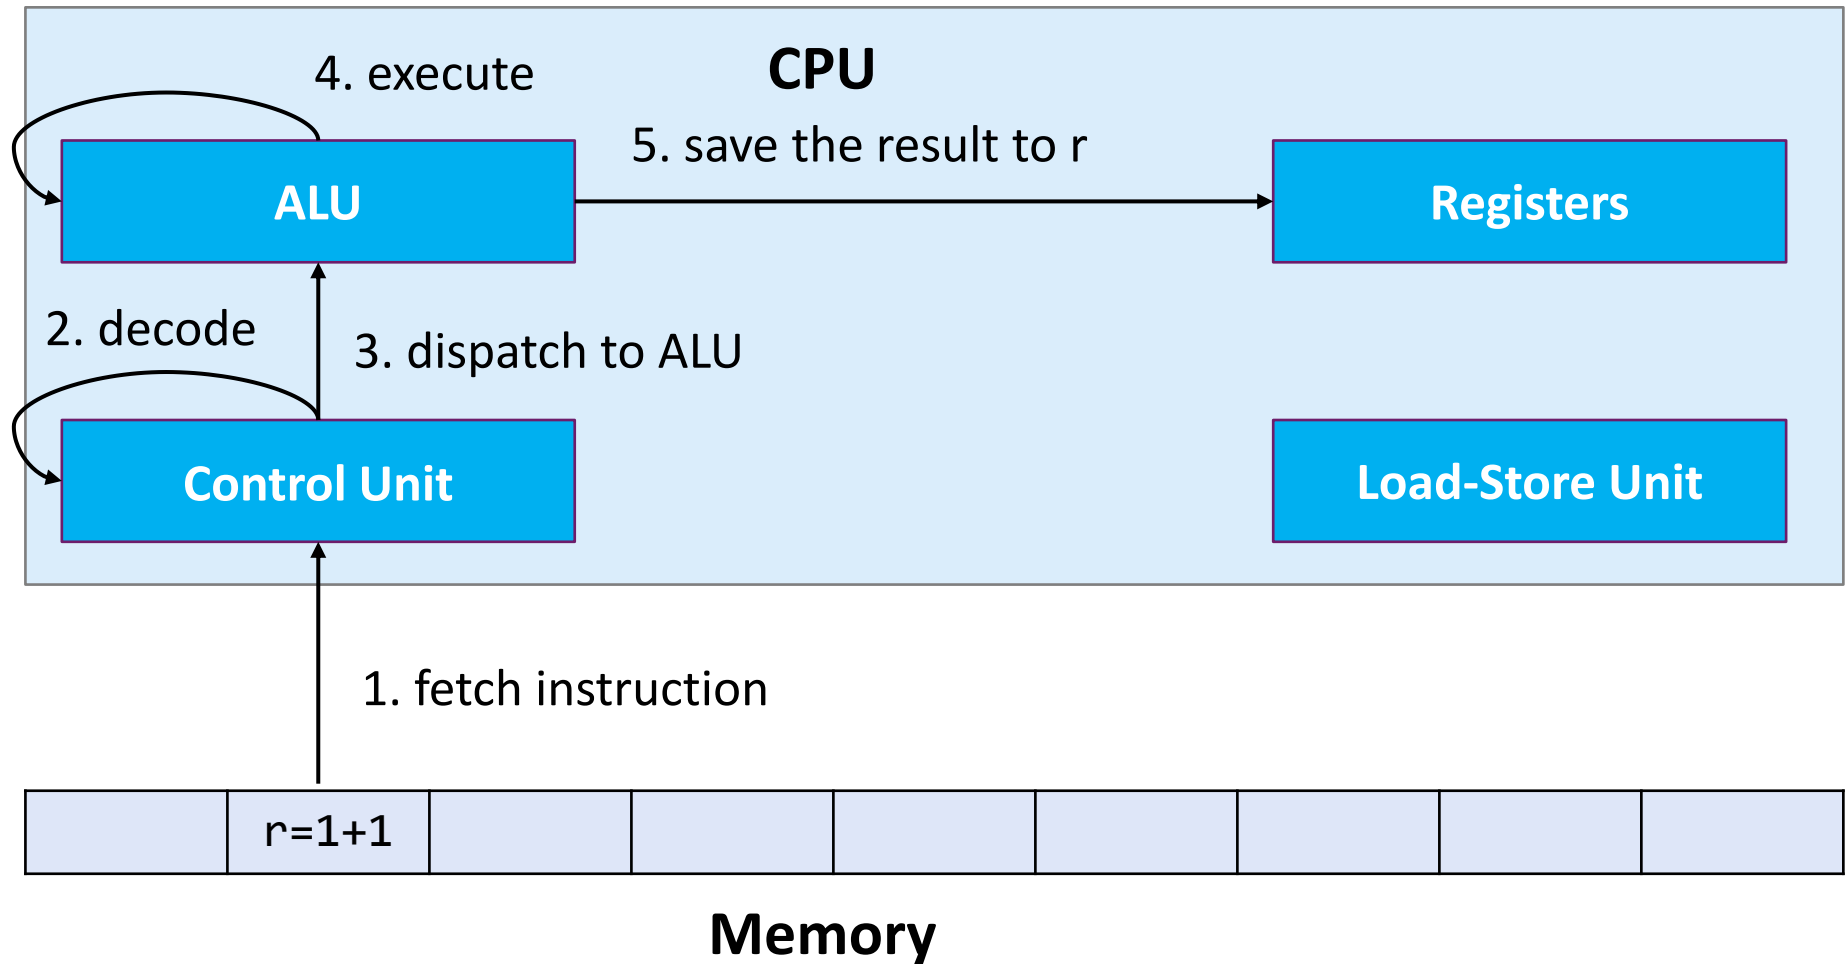


CPU: Central Processing Unit

- ❑ Arithmetic logic unit: Perform arithmetic/logic operations.
- ❑ Control unit: Fetch instructions from memory; decode and dispatch.
- ❑ Load-store unit: Load/store data from/to memory.
- ❑ Registers: Supply operands to the ALU and save the execution result.




Example: $r = 1 + 1$



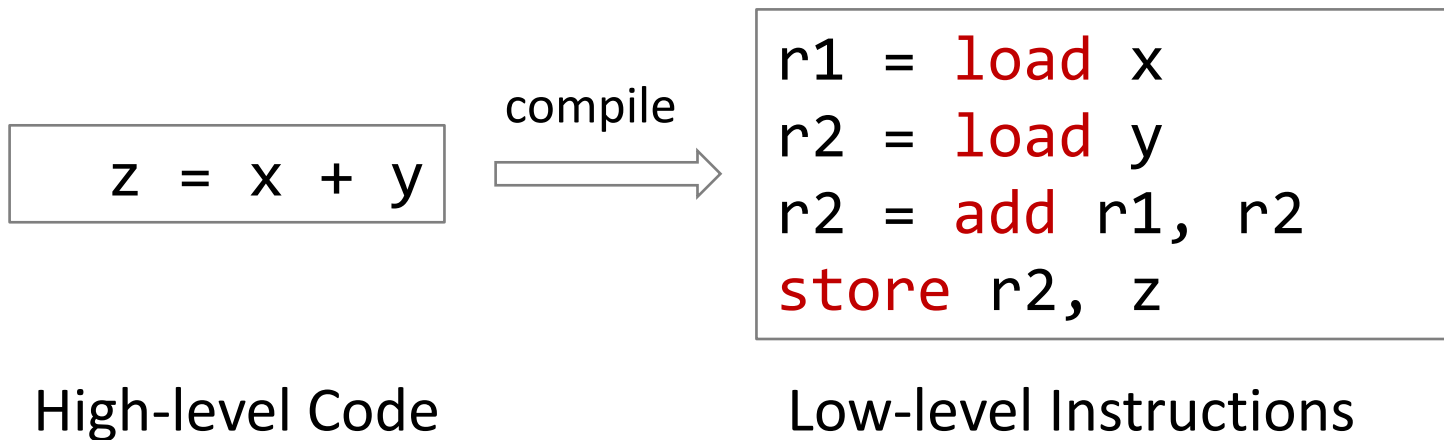
Memory (Random-Access Memory or RAM)

- ❑ Read or write a data item based on its address in memory.
- ❑ Each data unit generally takes 1 byte or 8 bits.
- ❑ Access any data with the same latency irrespective of their locations.



Example: $z = x + y$

- ❑ x and y are two integers stored in memory.
- ❑ Compiling the code generates four instructions.



Memory View

- ❑ We may assume each instruction is encoded as 32 bits (4 bytes).
- ❑ The four instructions are stored consecutively in memory with 16 bytes.
- ❑ The variables *x*, *y*, and *z* are symbolic names that will be replaced with their corresponding memory addresses.

	r1=load 0x18	r2=load 0x1B	r2=add r1,r2	store r2,0x20		x	y	z
--	-----------------	-----------------	-----------------	------------------	--	---	---	---

Address: 0x00 0x04 0x08 0x0B 0x10 0x14 0x18 0x1B 0x20

Memory

Demonstration with Linux

Write a simple C program (you may generate via AI)

```
#: gcc target.c -o target
```

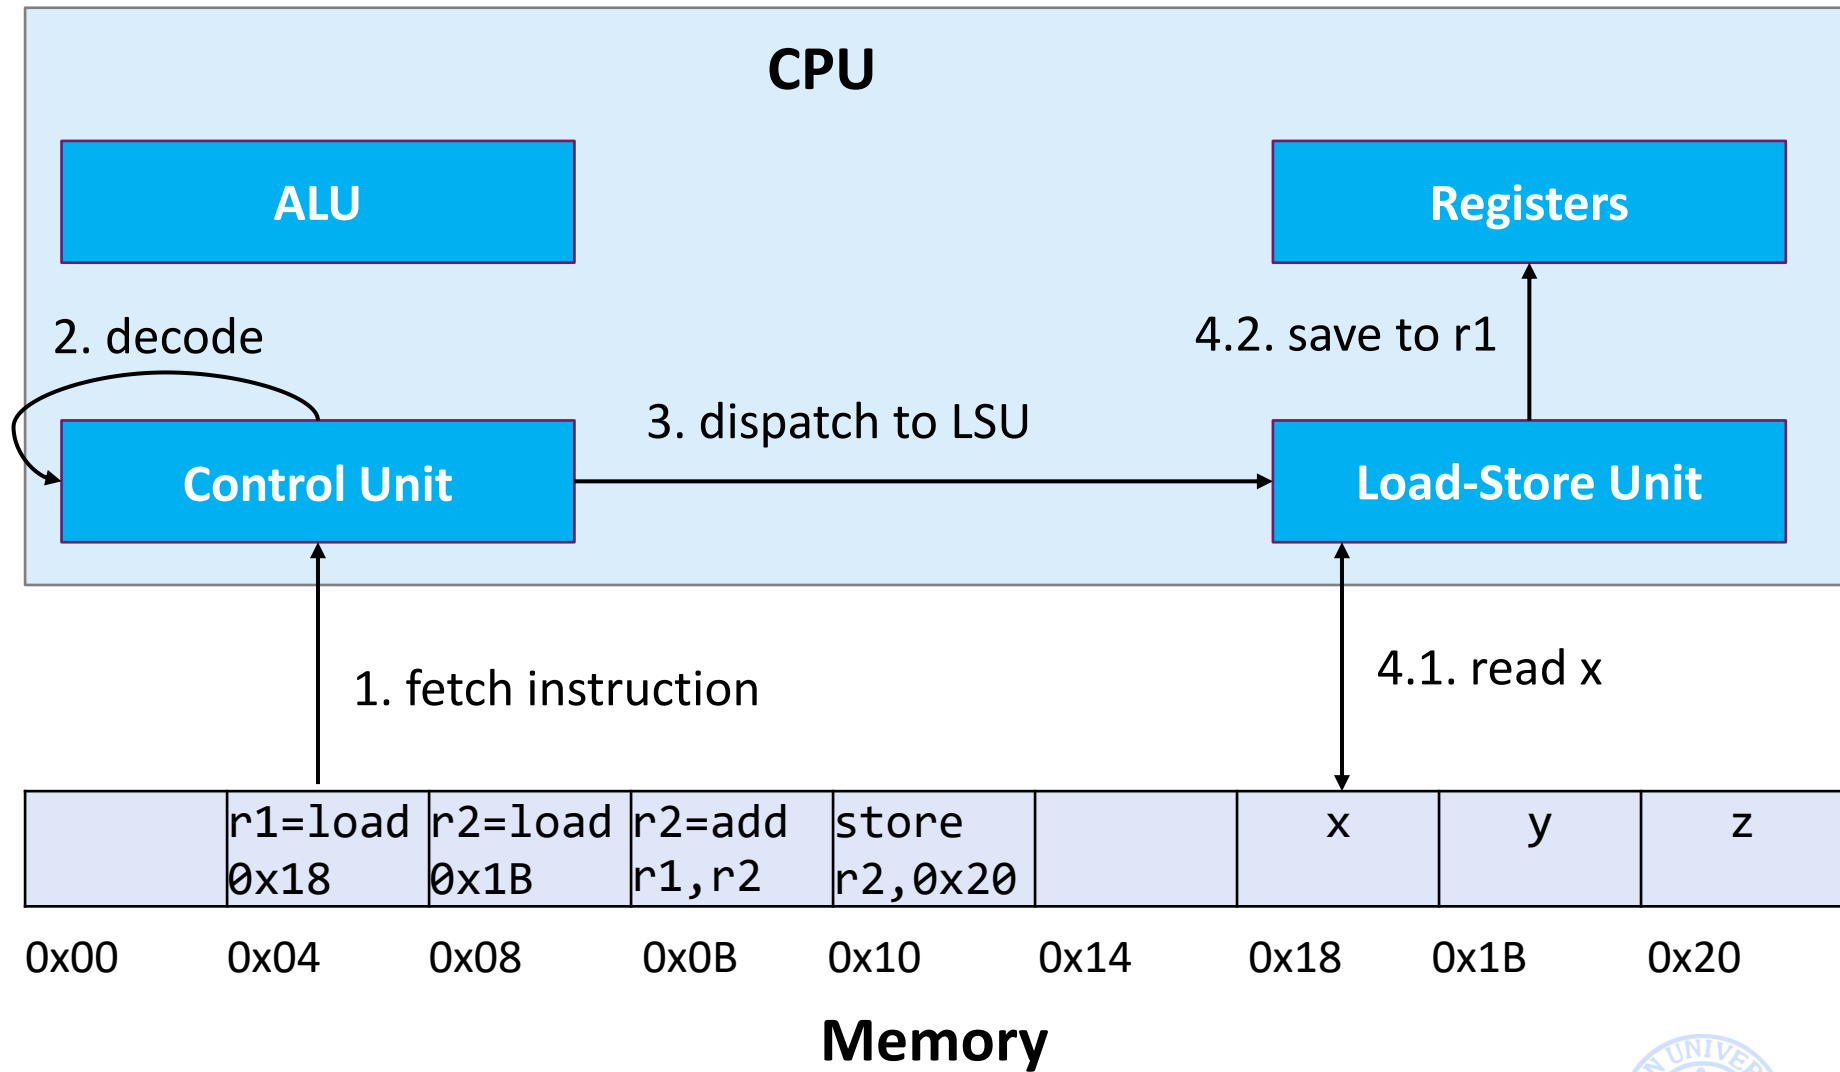
Single step execution via GDB

```
#: gdb target  
(gdb) disass main  
(gdb) break main  
(gdb) run  
(gdb) si
```

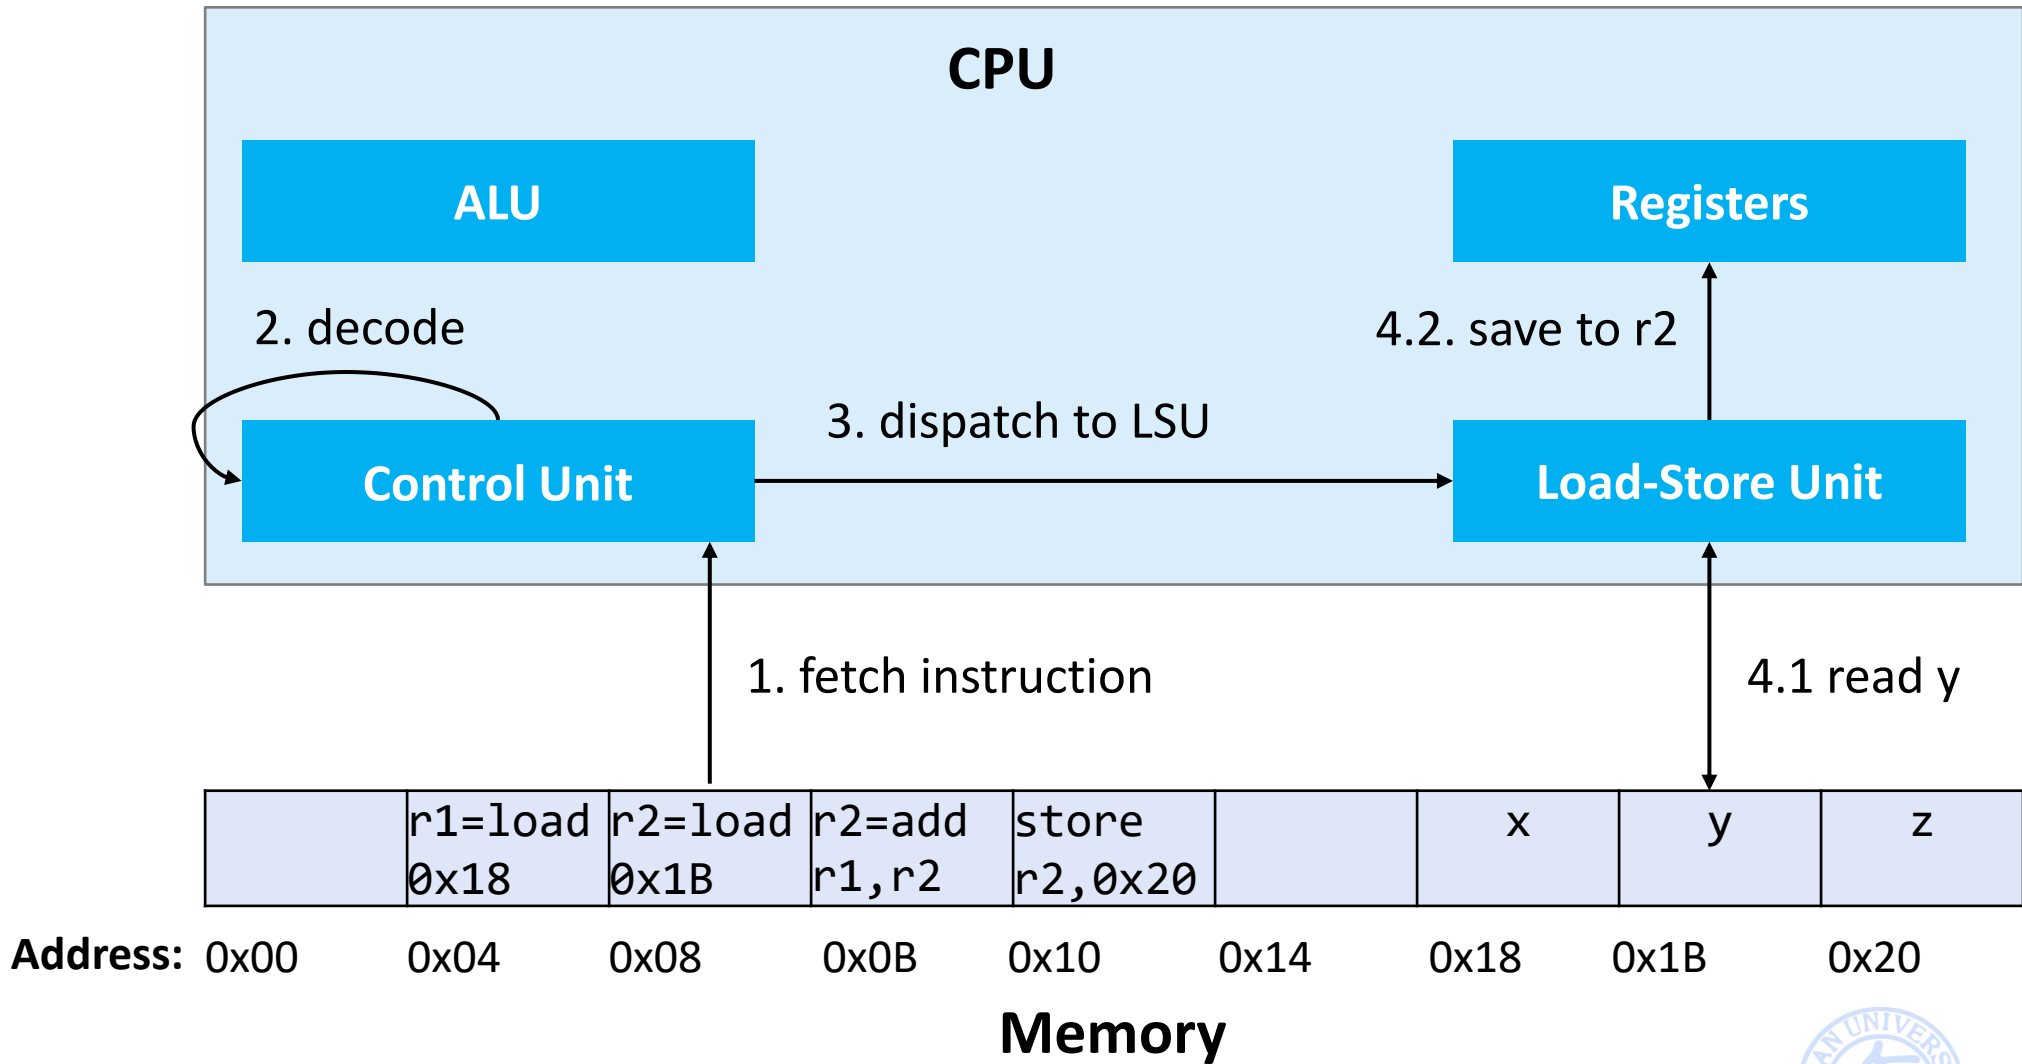
View the whole memory content of a process:

```
#: cat /proc/$pid/maps
```

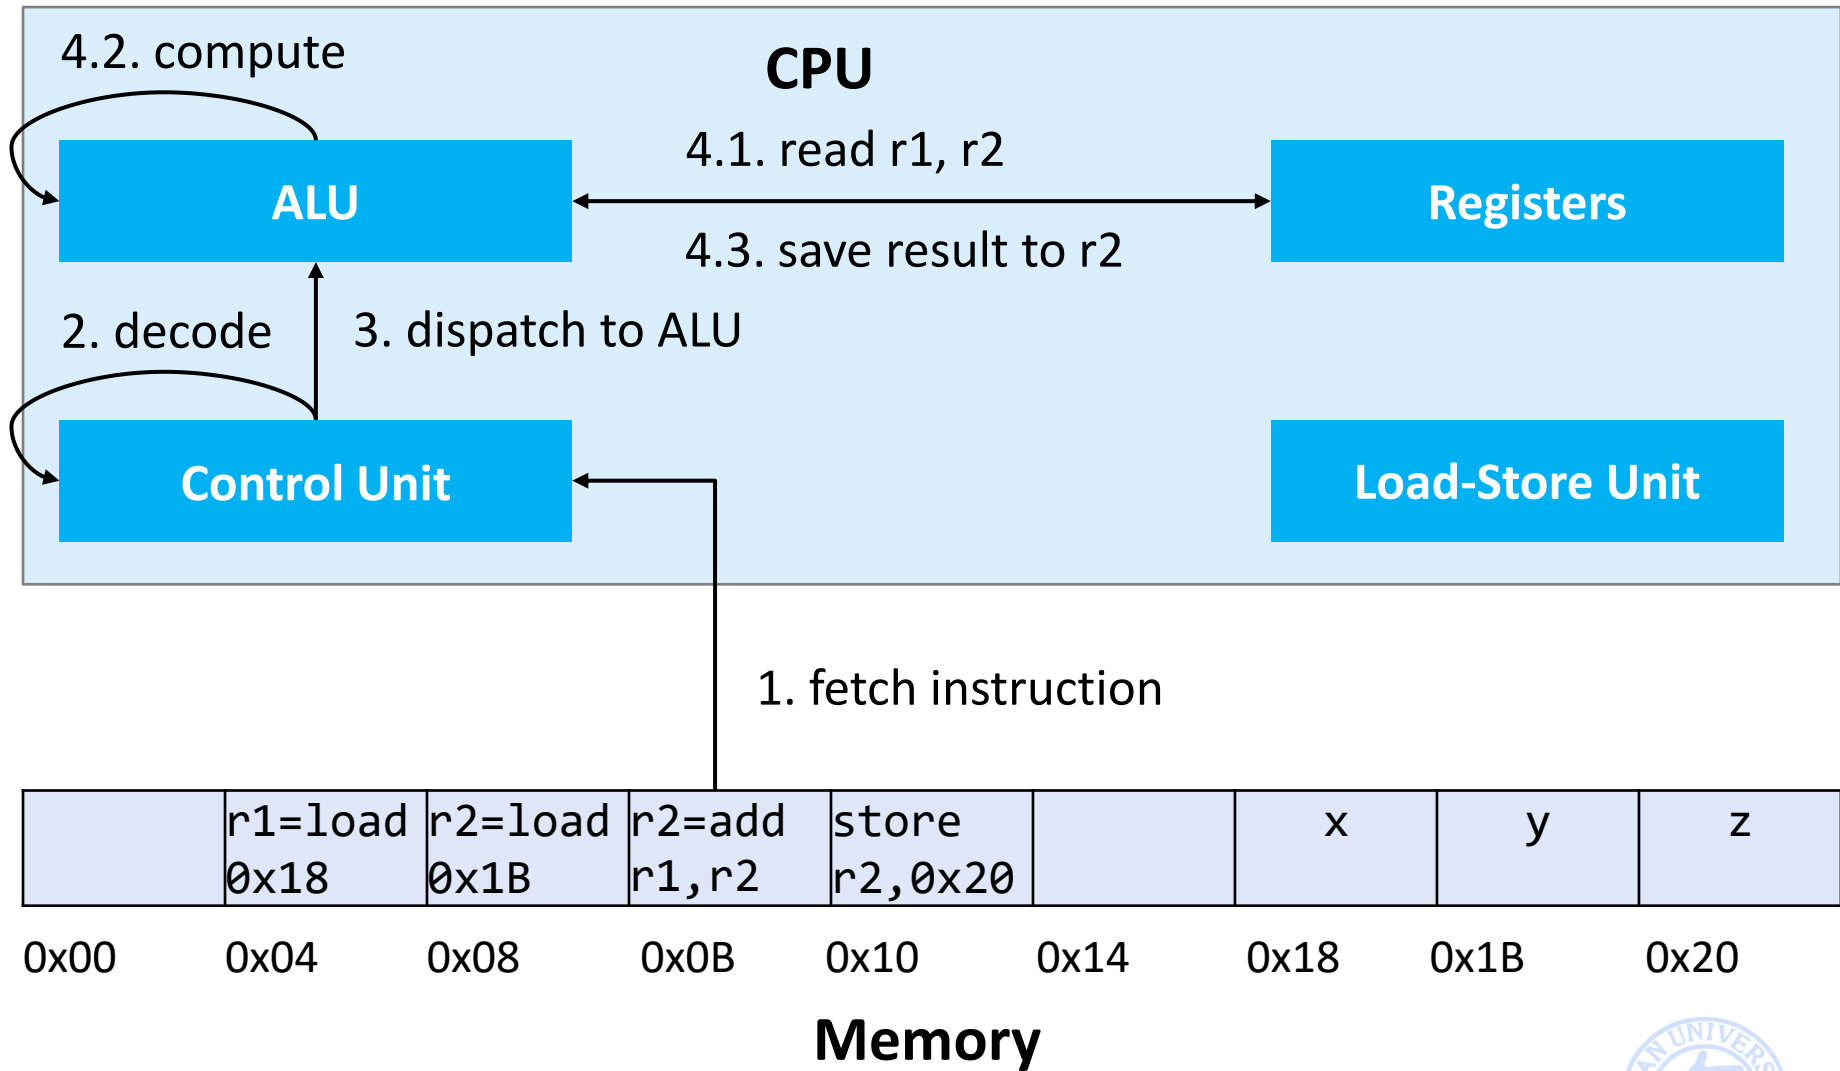
Step 1: r1 = load x



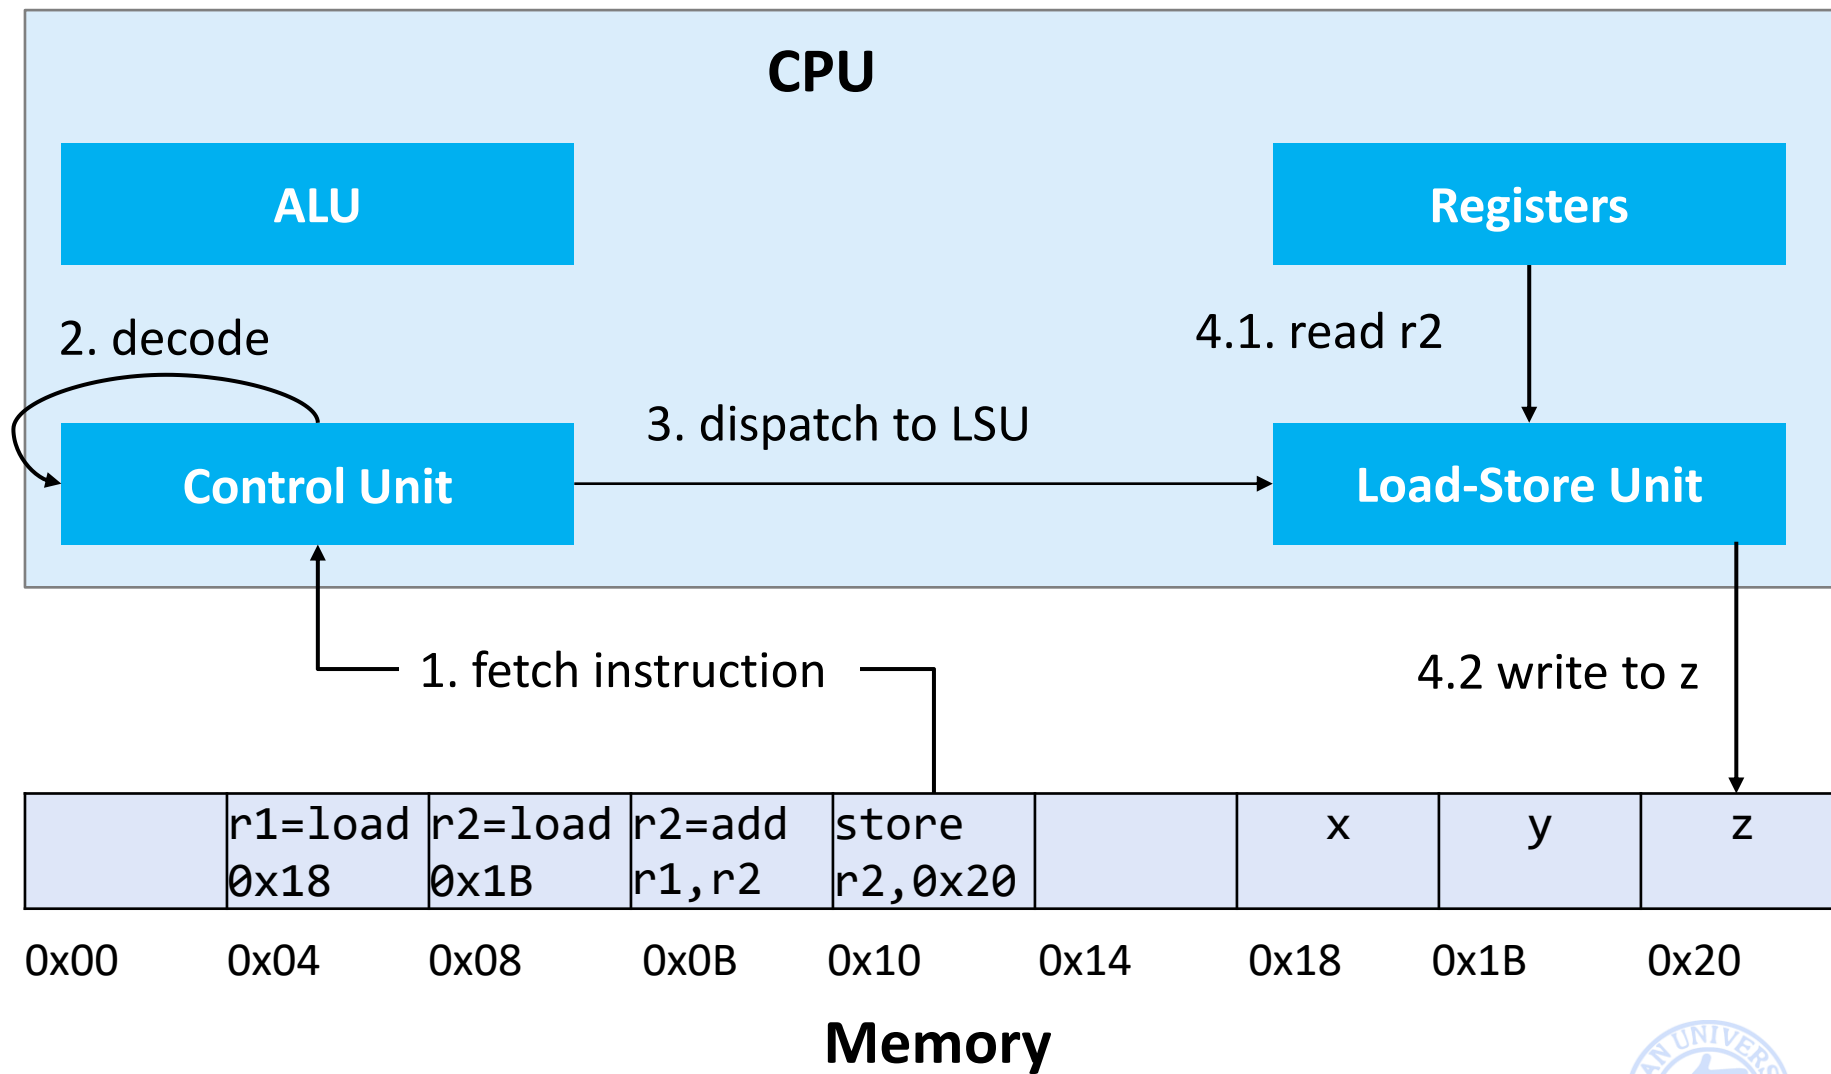
Step 2: r1 = load y



Step 3: r2=add r1, r2



Step 4: store r2, z



Sample CPU Core: Skylake

□ Frontend:

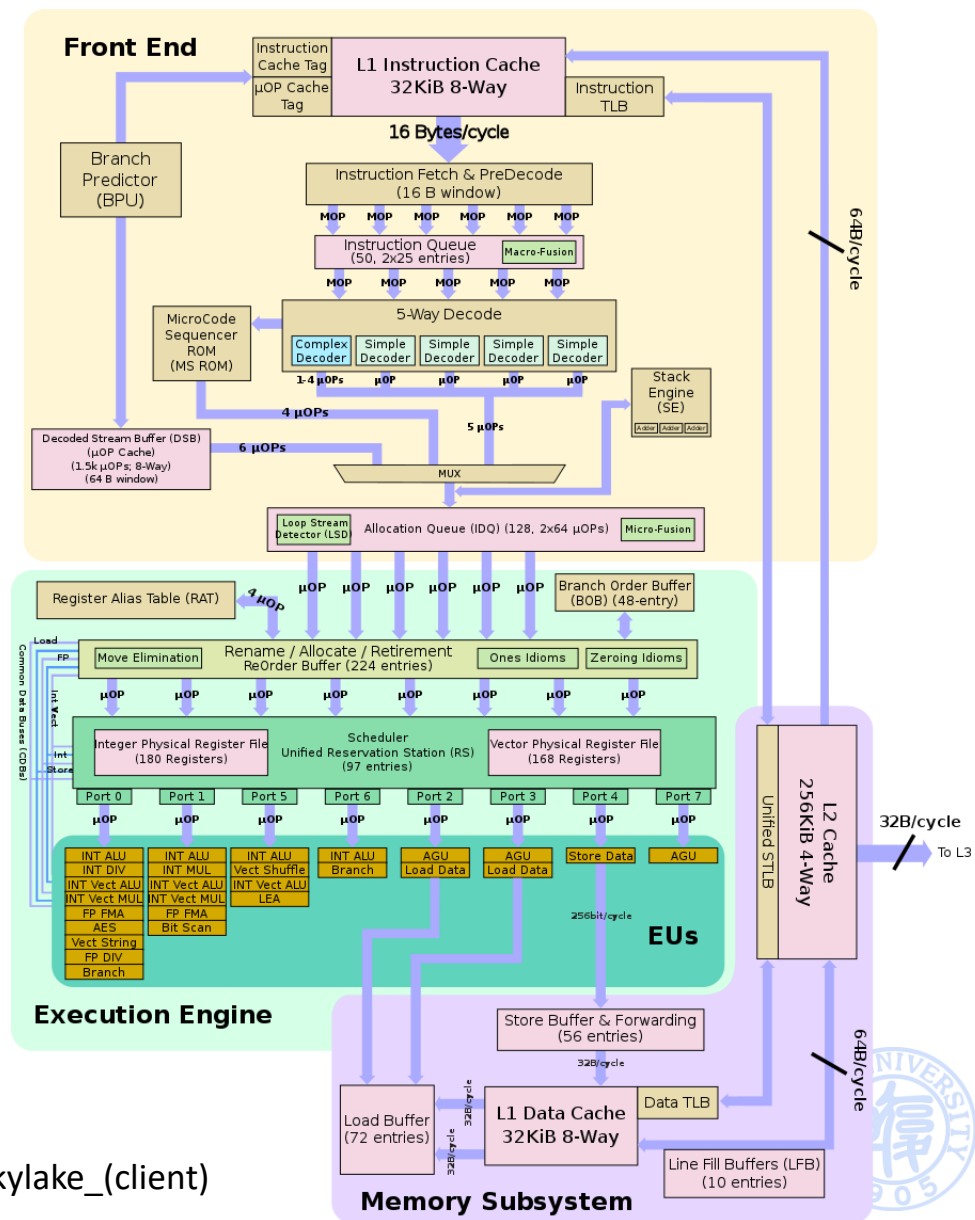
- Fetch, decode, ...

❑ Execution engine (backend):


- ALU, load-store unit, ...

❑ Memory subsystem

- Cache



Intel i7 with Four Cores



Demonstration (PC)

```
#: lscpu
Architecture:                x86_64
  CPU op-mode(s):            32-bit, 64-bit
  Address sizes:              39 bits physical, 48 bits virtual
  Byte Order:                 Little Endian
CPU(s):                       8
  On-line CPU(s) list:       0-7
Vendor ID:                    Genuine Intel
  Model name:                 Intel(R) Core(TM) i7-9700T CPU @ 2.00GHz
    CPU family:                6
    Model:                     158
    Thread(s) per core:        1
    Core(s) per socket:        8
    Socket(s):                 1
    Stepping:                  13
    CPU max MHz:               4300.0000
    CPU min MHz:               800.0000
...

```

Demonstration (Server)

```
#: lscpu
Architecture:                x86_64
  CPU op-mode(s):            32-bit, 64-bit
  Address sizes:             46 bits physical, 48 bits virtual
  Byte Order:                Little Endian
CPU(s):                      80
  On-line CPU(s) list:      0-79
Vendor ID:                   GenuineIntel
  Model name:                Intel(R) Xeon(R) Gold 6242R CPU @ 3.10GHz
    CPU family:              6
    Model:                   85
    Thread(s) per core:     2
    Core(s) per socket:     20
    Socket(s):               2
    Stepping:                7
    CPU max MHz:             4100.0000
    CPU min MHz:             1200.0000
...

```

CPU Frequency

❑ How many clock cycles the CPU completes per second.

- Shorter clock period => Higher CPU frequency

❑ Clock ticks synchronize all operations inside the CPU

- Clock cycle is determined by the critical path of the CPU's circuits.
- Signals travel from one register → combinational logic → next register in one cycle.
- Smaller transistor size => Faster signal propagation => Shorter critical path.
- Simple instruction take 1 cycle (*e.g.*, adding two registers).
- Complex instruction may take several cycles (*e.g.*, memory access, division).

Question

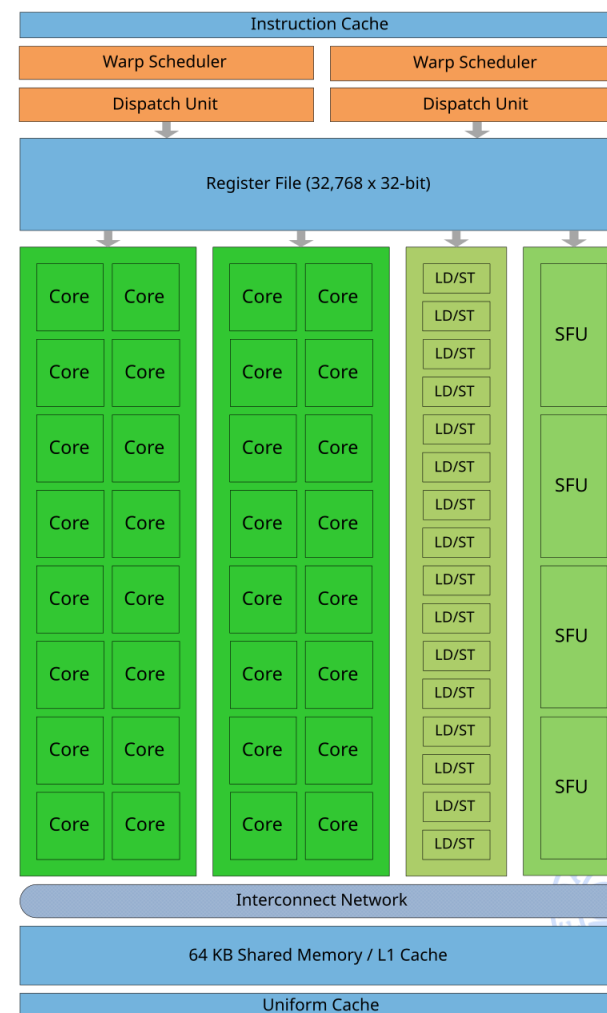


- ❑ If a CPU employs 32 bits to represent a memory address, what is the maximum memory space the CPU can support?

General-Purpose Computing with Graphics Processing Units

- ❑ GPU has thousands of cores and a highly parallel architecture.
- ❑ It was originally designed for graphics.
- ❑ GPGPU uses GPU to perform computation typically handled by CPU

Example GPU: Nvidia Fermi

- ❑ A GPU has multiple streaming multiprocessors (SMs).
- ❑ Each SM consists of 32 cores.



Parameters of Different NVidia GPUs

Arch	Product	#. of Cores		TFLOPS (FP32)	Memory	Bandwidth GB/s	Price
		CUDA	Tensor				
Pascal	GTX 1080	2560	-	8.9	8GB GDDR5X	352	\$599
Pascal	P100	3584	-	9.3	16GB HBM2	732	\$5,000+
Turing	RTX 2080	2944	46	10.1	8GB GDDR6	448	\$699
Ampere	RTX 3080	8704	328	29.8	10GB GDDR6X	912	\$699
Ampere	A100	6912	432	19.5	40/80GB HBM2e	1555	\$11,000
Hopper (PCI-e/SXM)	H100	16896	528	51	80GB HBM3	2039	\$30,000
Ada Lovelace	RTX 4090	16384	512	82.6	24GB GDDR6X	1008	\$1,599
Blackwell	RTX 5090	21760	680	104.8	32	1792	\$1,999

Demonstration

```
#: lspci | grep -i nvidia
```

```
01:00.0 VGA compatible controller: NVIDIA Corporation TU102 [GeForce RTX 2080 Ti Rev. A] (rev a1)
01:00.1 Audio device: NVIDIA Corporation TU102 High Definition Audio Controller (rev a1)
01:00.2 USB controller: NVIDIA Corporation TU102 USB 3.1 Host Controller (rev a1)
01:00.3 Serial bus controller [0c80]: NVIDIA Corporation TU102 USB Type-C UCSI Controller (rev a1)
02:00.0 VGA compatible controller: NVIDIA Corporation TU102 [GeForce RTX 2080 Ti Rev. A] (rev a1)
02:00.1 Audio device: NVIDIA Corporation TU102 High Definition Audio Controller (rev a1)
02:00.2 USB controller: NVIDIA Corporation TU102 USB 3.1 Host Controller (rev a1)
02:00.3 Serial bus controller [0c80]: NVIDIA Corporation TU102 USB Type-C UCSI Controller (rev a1)
```

```
#: nvidia-smi
```

```
Mon Sep 15 11:16:32 2025
```

-----+-----									
NVIDIA-SMI 525.89.02				Driver Version: 525.89.02			CUDA Version: 12.0		
-----+-----									
GPU	Name	Persistence-M			Bus-Id	Disp.A	Volatile	Uncorr. ECC	
Fan	Temp	Perf	Pwr:Usage/Cap			Memory-Usage		GPU-Util	Compute M.
=====+=====									
0	NVIDIA	GeForce ...	Off		00000000:01:00.0	Off			N/A
27%	44C	P0	58W / 250W			0MiB / 11264MiB		0%	Default
-----+-----									
1	NVIDIA	GeForce ...	Off		00000000:02:00.0	Off			N/A
32%	45C	P0	31W / 250W			0MiB / 11264MiB		0%	Default
-----+-----									



2. ISA and Assembly Code

ISA: Instruction Set Architecture

- ❑ **ISA is the interface that defines what a processor can do and how software communicates with the hardware.**
- ❑ **RISC: Reduced Instruction Set Computer**
 - Separate data load/store and computation into different instructions
 - *e.g.*, AArch/ARM、 RISC-V
- ❑ **CISC: Complex Instruction Set Computer**
 - An instruction can do both data load/store and computation
 - *e.g.*, X86, X86-64

ARMv8-A Instructions

❑ The operands of instructions can be registers, instant numbers, *etc.*

❑ In the following example:

➤ x0..x5 are registers

➤ #5, #10 are instant numbers

Instruction

Semantic

```
mov x0, #5
```

$x0 = 5$

```
mov x1, #10
```

$x1 = 10$

```
add x2, x0, x1
```

$x2 = x0 + x1$

```
sub x3, x0, x1
```

$x3 = x0 - x1$

```
mul x4, x0, x1
```

$x4 = x0 * x1$

```
sdiv x5, x0, x1
```

$x5 = x0 / x1$

ARMv8-A Instructions: Data Load/Store

❑ Memory address as an operand:

➤ $[x]$ indicates the memory address pointed by x

Instruction	Semantic
-------------	----------

<code>ldr x1, [x0]</code>	$x1 = [x0]$
---------------------------	-------------

<code>ldr x2, [x0, #12]</code>	$x2 = [x0 + 12]$
--------------------------------	------------------

<code>add x3, x1, x2</code>	$x3 = x1 + x2,$
-----------------------------	-----------------

<code>str x3, [x0, #8]</code>	$[x0 + 8] = x3$
-------------------------------	-----------------

<code>add x2, x1, [x0]</code>	Not allowed (CISC operation)
-------------------------------	------------------------------

ARMv8-A Instructions: Control

❑ Control instructions change the normal flow of execution.

➤ *e.g.*, jumping to another part of code.

address Instructions

0x20	cmp	x0, x1
0x24	ble	.addr1
...	...	
0x104	.addr1:	
...	...	

[flag] = compare(x0, x1)

if(flag:<=) PC = 0x104

...

PC is a register that always holds the memory address of the next instruction.

Instructions in Memory

□ 32bit fixed-length for ARMv8-A.

□ Limitation: Cannot employ some large numbers as operands.

	r1= load x	r2= load y	r2=add r1, r2	store r2, z		x	y	z
Address: 0x00	0x04	0x08	0x0B	0x10	0x14	0x18	0x1B	0x20

31	30	29	28	27	26	25	24	23	22	21	20	16	15	10	9	5	4	0
sf	0	0	0	1	0	1	1	shift	0	Rm			imm6			Rn		Rd

op S

add rn, rm, rd

31	30	29	28	27	26	25	24	23	22	21	10	9	5	4	0
sf	0	0	1	0	0	0	1	0	sh	imm12					Rd

op S

add rd, rn, imm12
add rd, rn, #4095
add rd, rn, #4097

Invalid, because 4097 needs 13 bits.



Sample Program

❑ The following code computes the largest Fibonacci number less than 100 and saves the result to the register x0 (aarch64) or eax (x86).

➤ 1+1=2, 1+2=3, 2+3=5, 3+5=8, 5+8=13, 8+13=21, 13+21=34, 21+34=55, 34+55=89

```
_main:
    mov x0, #1
    mov x2, #1

_loop:
    mov x1, x0
    mov x0, x2
    add x2, x2, x1
    cmp x2, #100
    ble _loop
    ret
```

AArch64 Version

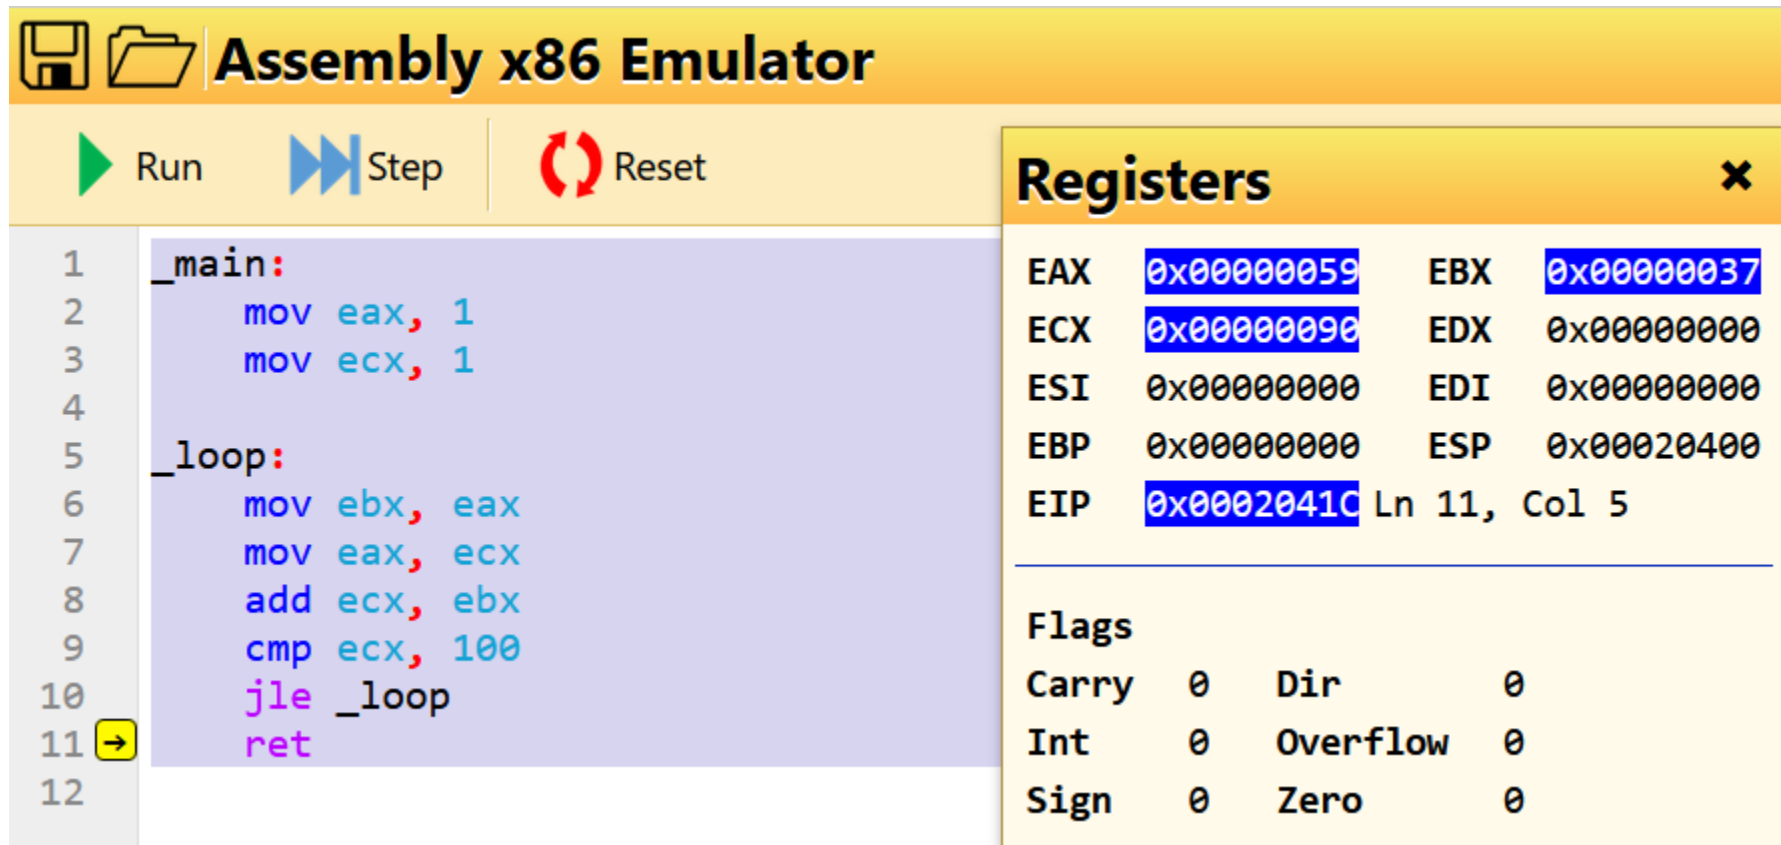
```
_main:
    mov eax, 1
    mov ecx, 1

_loop:
    mov ebx, eax
    mov eax, ecx
    add ecx, ebx
    cmp ecx, 100
    jle _loop
    ret
```

X86 Version

ecx = ecx + ebx

Examine The Code via an Emulator



The screenshot displays the 'Assembly x86 Emulator' window. The assembly code is as follows:

```
1  _main:
2      mov  eax, 1
3      mov  ecx, 1
4
5  _loop:
6      mov  ebx, eax
7      mov  eax, ecx
8      add  ecx, ebx
9      cmp  ecx, 100
10     jle  _loop
11     ret
12
```

The 'Registers' panel on the right shows the following values:

Register	Value
EAX	0x00000059
ECX	0x00000090
ESI	0x00000000
EBP	0x00000000
EIP	0x0002041C
EBX	0x00000037
EDX	0x00000000
EDI	0x00000000
ESP	0x00020400

The 'Flags' panel shows the following status:

Flag	Value
Carry	0
Int	0
Sign	0
Dir	0
Overflow	0
Zero	0

The EIP register is highlighted with a blue box, and the text 'Ln 11, Col 5' is displayed next to it.

Exercise

□ Modify the previous code to compute the largest accumulated value less than 100, starting from 1.

➤ $\text{result} = \text{Max}(1+2+3+4+5+\dots)$, s.t. $\text{result} < 100$


3. I/O and More

Computer Architecture

❑ Besides CPU and memory, there are peripheral devices

- Input device: keyboard, mouse, microphone, *etc.*
- Output device: display, printer, *etc.*
- Storage/Network: hard disk, USB flash drive, *etc.*

❑ How does the CPU read/write peripheral devices?

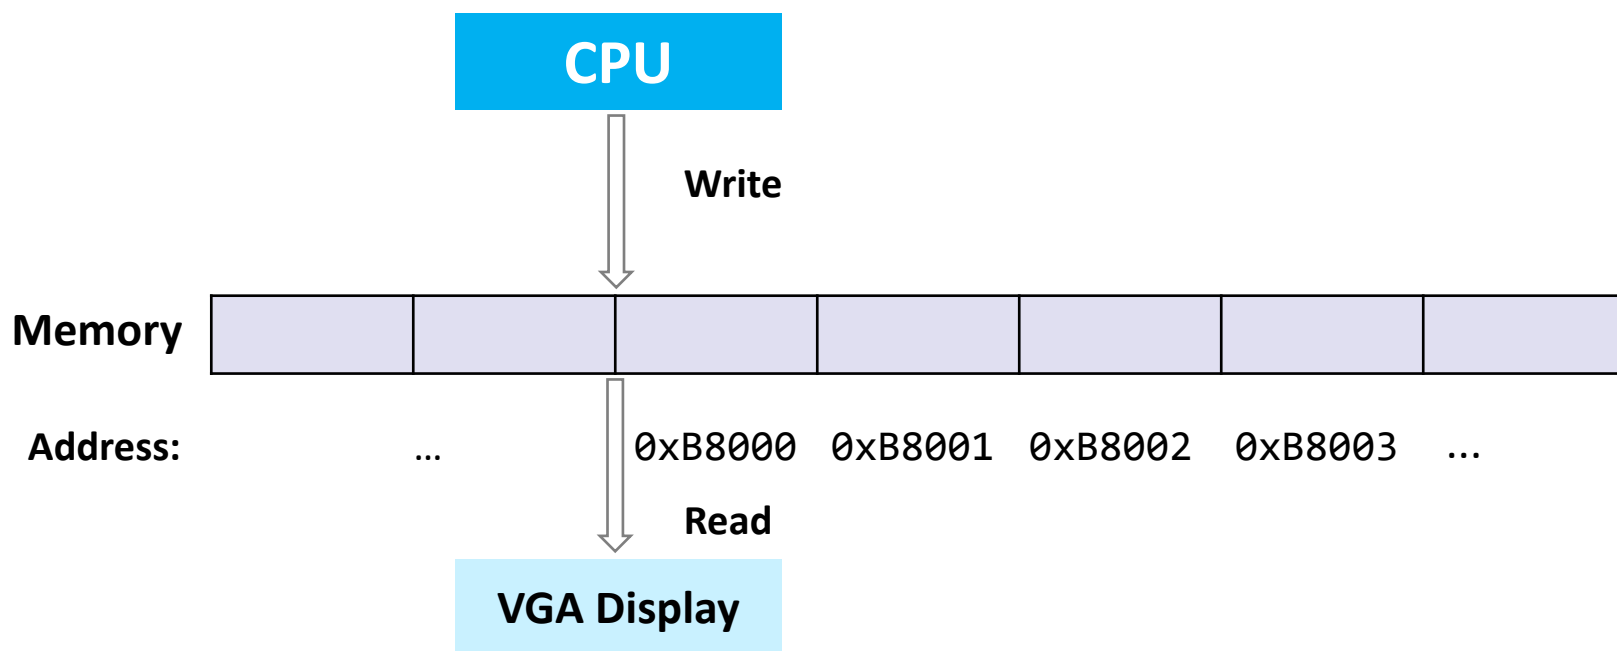


control: 
data: 



Memory-Mapped I/O

- ❑ Read/write data to I/O devices via memory access.
- ❑ Addressing achieved via PCI (peripheral component interconnect).




Output: Video as an Example

- ❑ VGA (video graphics array): video display controller 15-pin.
- ❑ Mapped to the PC's memory: 0xA0000 - 0xBFFFF.
- ❑ 0xB8000 for color text mode (32 KB).

```
mov ax, 0xB8000    ; 0xB8000 is the VGA address
mov es, ax         ; move the content of ax to es
mov di, 0          ; move the value 0 to di register
mov al, 'A'        ; move the value 'A' to al register
mov ah, 0x0F       ; move the value 0x0F (color) to ah
mov [es:di], al    ; move the content of al to memory address:
                  ; es+di
mov [es:di+1], ah  ; move the content of ah to memory address:
                  ; es+di+1
```

Alternative Approach: Port-based I/O

- ❑ Read/write data to I/O devices via port.
- ❑ For example, the PS/2 interface employs two ports:
 - 0x60 for command
 - 0x64 for data



Handling Input is More Complicated

❑ When/how many data arrives is not controlled by the CPU.

- Polling: let the CPU check particular signals periodically? How frequently?
- By interrupt (trigger-based) via particular CPU pins

Reading the keyboard via polling

```
wait_for_input:
    in al, 0x64    ; Read the signal from port 0x64 into AL
    test al, 0x01 ; Test whether bit 0 of AL is set (if data is ready)
    jz wait_for_input ; If zero, keep polling until input is available

read_input:
    in al, 0x60    ; Read the character from port 0x60
```

Reading the keyboard via interrupts

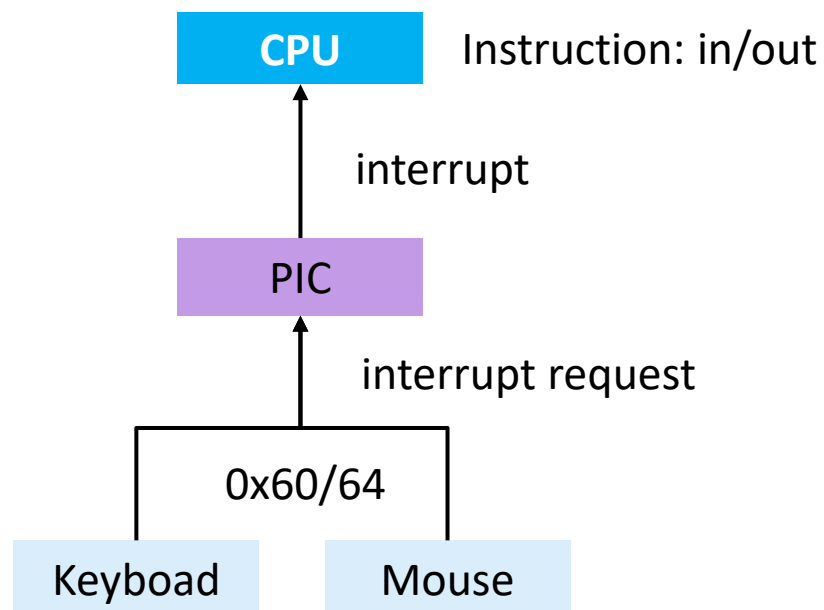
```
in al, 0x60 ; ; Read the byte from port 0x60
...
out 0x20, 0x20 ; Send End-Of-Interrupt command to port 0x20
```



Input is More Complicated

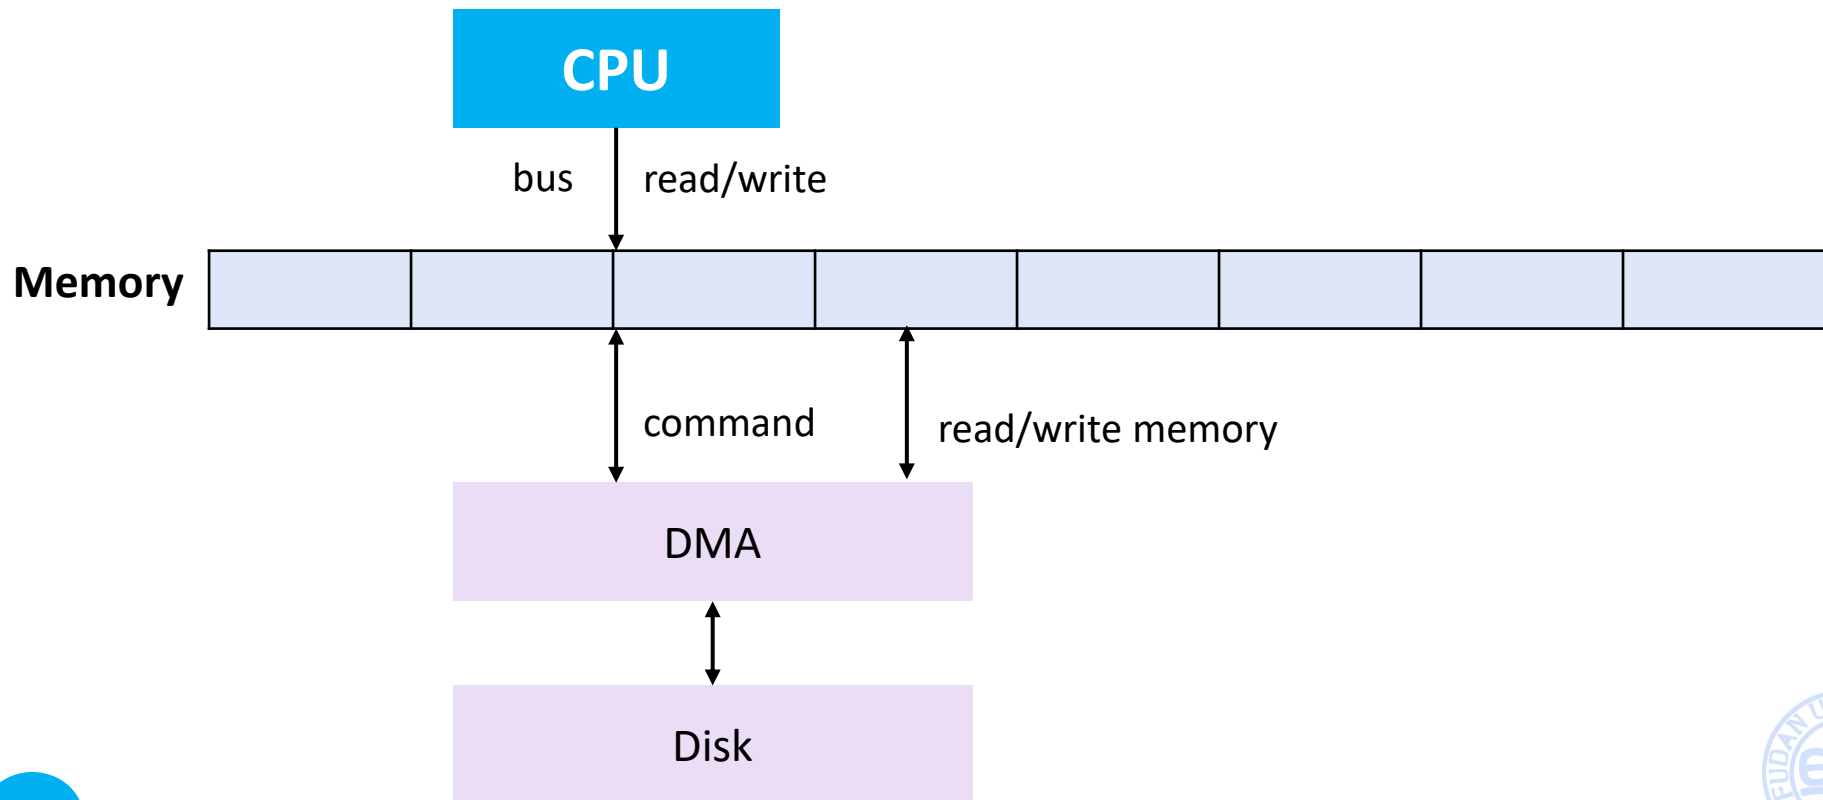
❑ What if data arrives from multiple input devices simultaneously?

- First come first serve? May lead to denial-of-service
- Priority/queuing via PIC (programmable interrupt controller)



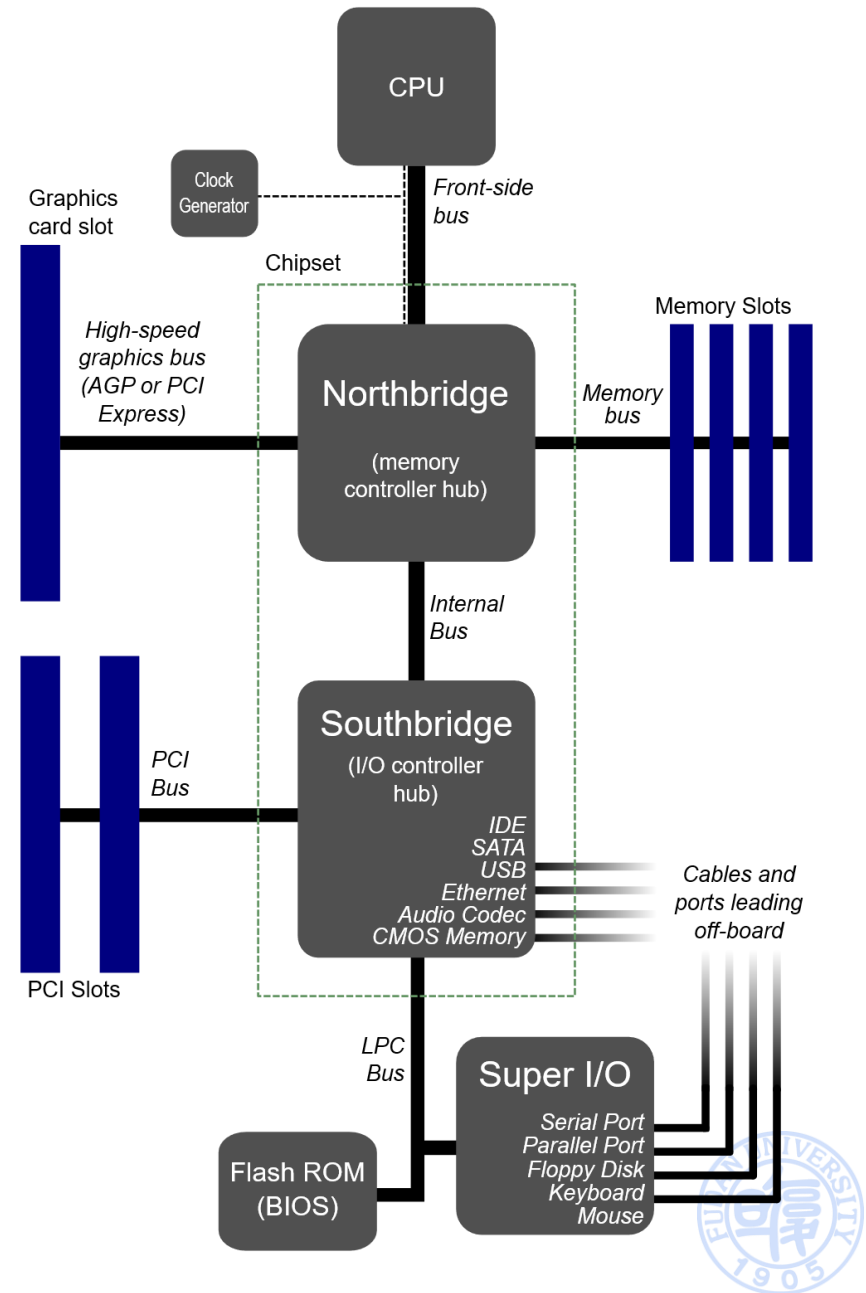
Disk I/O

- ❑ Disk read/write operations are much slower than memory.
- ❑ Offload the workload to disk controller via direct memory access.




Motherboard

- ❑ Northbridge: handle high-performance tasks, connected directly to a CPU.
- ❑ Southbridge: handle low speed tasks, connected via the northbridge



Motherboard



Demonstration

```
#: lspci
```

```
00:00.0 Host bridge: Intel Corporation 8th/9th Gen Core 8-core Desktop Processor Host Bridge/DRAM  
Registers [Coffee Lake S] (rev 0d)  
00:02.0 VGA compatible controller: Intel Corporation CoffeeLake-S GT2 [UHD Graphics 630] (rev 02)  
00:08.0 System peripheral: Intel Corporation Xeon E3-1200 v5/v6 / E3-1500 v5 / 6th/7th/8th Gen Core  
Processor Gaussian Mixture Model  
00:12.0 Signal processing controller: Intel Corporation Cannon Lake PCH Thermal Controller (rev 10)  
00:14.0 USB controller: Intel Corporation Cannon Lake PCH USB 3.1 xHCI Host Controller (rev 10)  
00:14.2 RAM memory: Intel Corporation Cannon Lake PCH Shared SRAM (rev 10)  
00:14.3 Network controller: Intel Corporation Cannon Lake PCH CNVi WiFi (rev 10)  
00:15.0 Serial bus controller: Intel Corporation Cannon Lake PCH Serial IO I2C Controller #0 (rev 10)  
00:16.0 Communication controller: Intel Corporation Cannon Lake PCH HECI Controller (rev 10)  
00:17.0 SATA controller: Intel Corporation Cannon Lake PCH SATA AHCI Controller (rev 10)  
00:1b.0 PCI bridge: Intel Corporation Cannon Lake PCH PCI Express Root Port #17 (rev f0)  
00:1f.0 ISA bridge: Intel Corporation Q370 Chipset LPC/eSPI Controller (rev 10)  
00:1f.3 Audio device: Intel Corporation Cannon Lake PCH cAVS (rev 10)  
00:1f.4 SMBus: Intel Corporation Cannon Lake PCH SMBus Controller (rev 10)  
00:1f.5 Serial bus controller: Intel Corporation Cannon Lake PCH SPI Controller (rev 10)  
00:1f.6 Ethernet controller: Intel Corporation Ethernet Connection (7) I219-LM (rev 10)  
01:00.0 Non-Volatile memory controller: Phison Electronics Corporation PS5013 E13 NVMe Controller (rev 01)
```

Demonstration

```
#: sudo lspci -vv
```

```
00:00.0 Host bridge: Intel Corporation 8th/9th Gen Core 8-core Desktop Processor Host Bridge/DRAM Registers
[Coffee Lake S] (rev 0d)
  DeviceName: Onboard - Other
  Subsystem: Dell 8th/9th Gen Core 8-core Desktop Processor Host Bridge/DRAM Registers [Coffee Lake S]
  Control: I/O- Mem+ BusMaster+ SpecCycle- MemWINV- VGASnoop- ParErr- Stepping- SERR+ FastB2B- DisINTx-
  Status: Cap+ 66MHz- UDF- FastB2B+ ParErr- DEVSEL=fast >TAbort- <TAbort- <MAbort+ >SERR- <PERR- INTx-
  Latency: 0
  Capabilities: [e0] Vendor Specific Information: Len=10 <?>
  Kernel driver in use: skl_uncore
  Kernel modules: ie31200_edac

00:02.0 VGA compatible controller: Intel Corporation CoffeeLake-S GT2 [UHD Graphics 630] (rev 02) (prog-if 00 [VGA
controller])
  DeviceName: Onboard - Video
  Subsystem: Dell CoffeeLake-S GT2 [UHD Graphics 630]
  Control: I/O+ Mem+ BusMaster+ SpecCycle- MemWINV- VGASnoop- ParErr- Stepping- SERR- FastB2B- DisINTx+
  Status: Cap+ 66MHz- UDF- FastB2B- ParErr- DEVSEL=fast >TAbort- <TAbort- <MAbort- >SERR- <PERR- INTx-
  Latency: 0, Cache Line Size: 64 bytes
  Interrupt: pin A routed to IRQ 142
  Region 0: Memory at 90000000 (64-bit, non-prefetchable) [size=16M]
  Region 2: Memory at 80000000 (64-bit, prefetchable) [size=256M]
  Region 4: I/O ports at 3000 [size=64]
  Expansion ROM at 000c0000 [virtual] [disabled] [size=128K]
  Capabilities: [40] Vendor Specific Information: Len=0c <?>

...
```

Several Concepts

- ❑ **Hardware:** Physical components of a computer system.
- ❑ **Firmware:** Specialized software embedded in hardware devices to control or manage the hardware directly.
 - *e.g.*, BIOS (basic input/output system) for interrupt handling.
- ❑ **Software:** Programs that run on the hardware.
 - *e.g.*, operating systems and applications.

Role of BIOS / UEFI

- ❑ BIOS/UEFI runs before the OS, during system boot.
- ❑ Assigns and initializes port addresses for some devices.
- ❑ Scans devices (PCI, ACPI, etc.) and assigns MMIO ranges and writes assigned addresses into device registers.

Role of the Operating System

❑ **After BIOS/UEFI initializes hardware, the OS takes over.**

➤ Talking to the hardware directly is difficult, error-prone, and hardware-specific.

❑ **OS manages hardware resources and provides services for apps.**

➤ Applications cannot talk to the hardware directly, ensuring safe device operations.



Summary

- ❑ A CPU consists of a control unit, ALU, LSU, and registers.
- ❑ Each CPU implements an ISA, programmed using assembly instructions.
- ❑ CPU communicates with peripheral devices via I/O ports, memory-mapped I/O, and interrupts.
- ❑ OS manages hardware and simplifies the usage of computer.