

Lecture 2: Computer Architecture

Hui Xu

xuh@fudan.edu.cn



Outline

- ❖ 1. CPU and ISA
- ❖ 2. I/O and More
- ❖ 3. In-class Practice

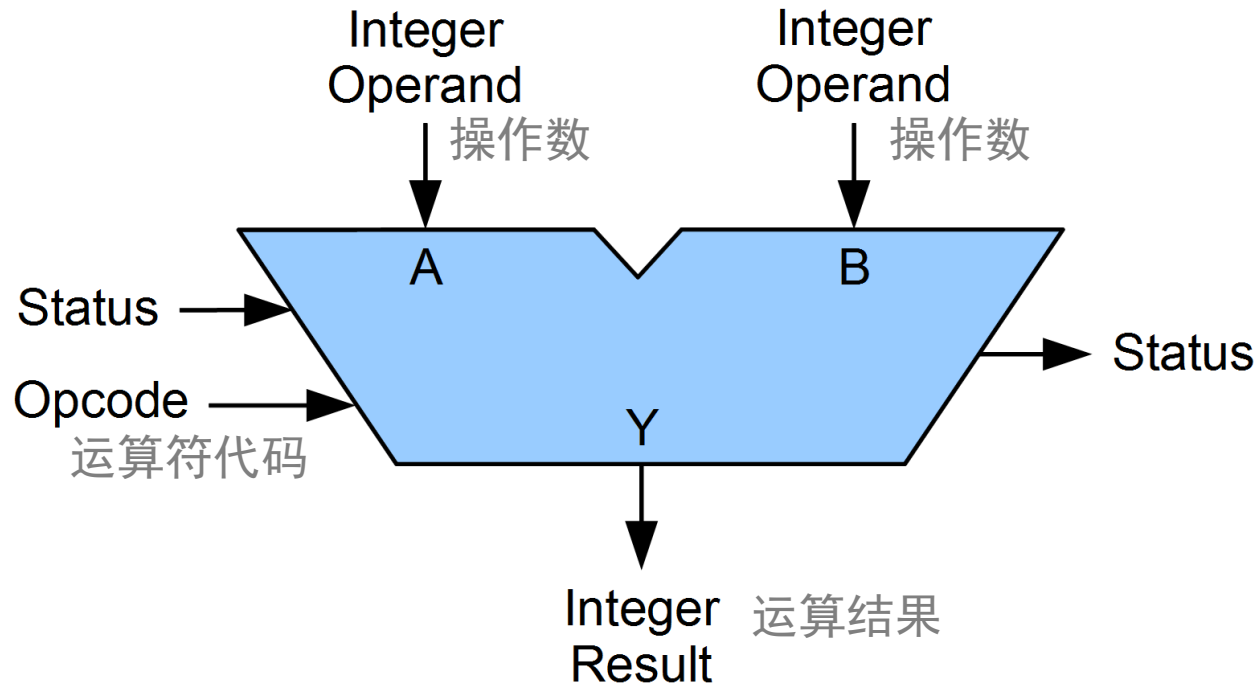
1. CPU and ISA

Recall: Transistor => Logic Gate => ALU

晶体管

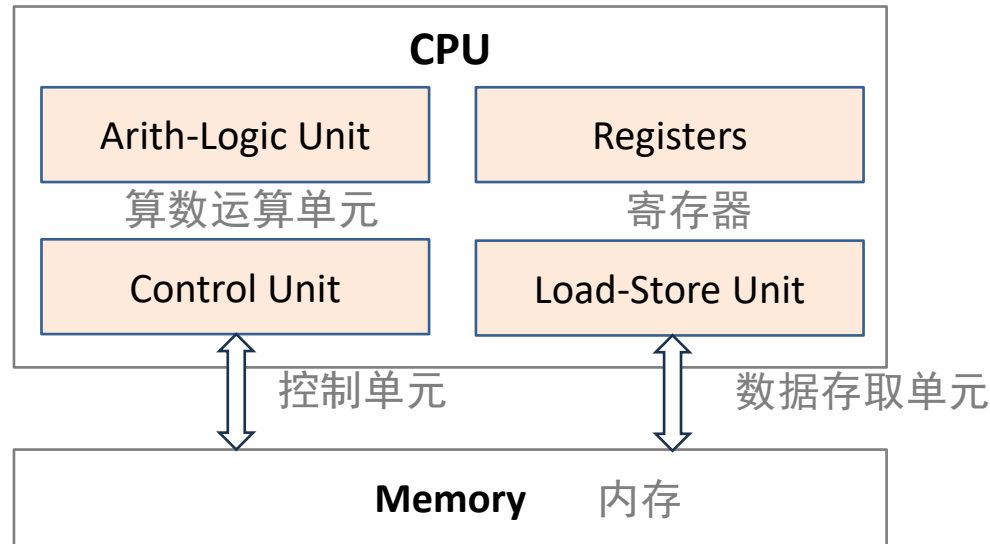
逻辑门

算数运算单元



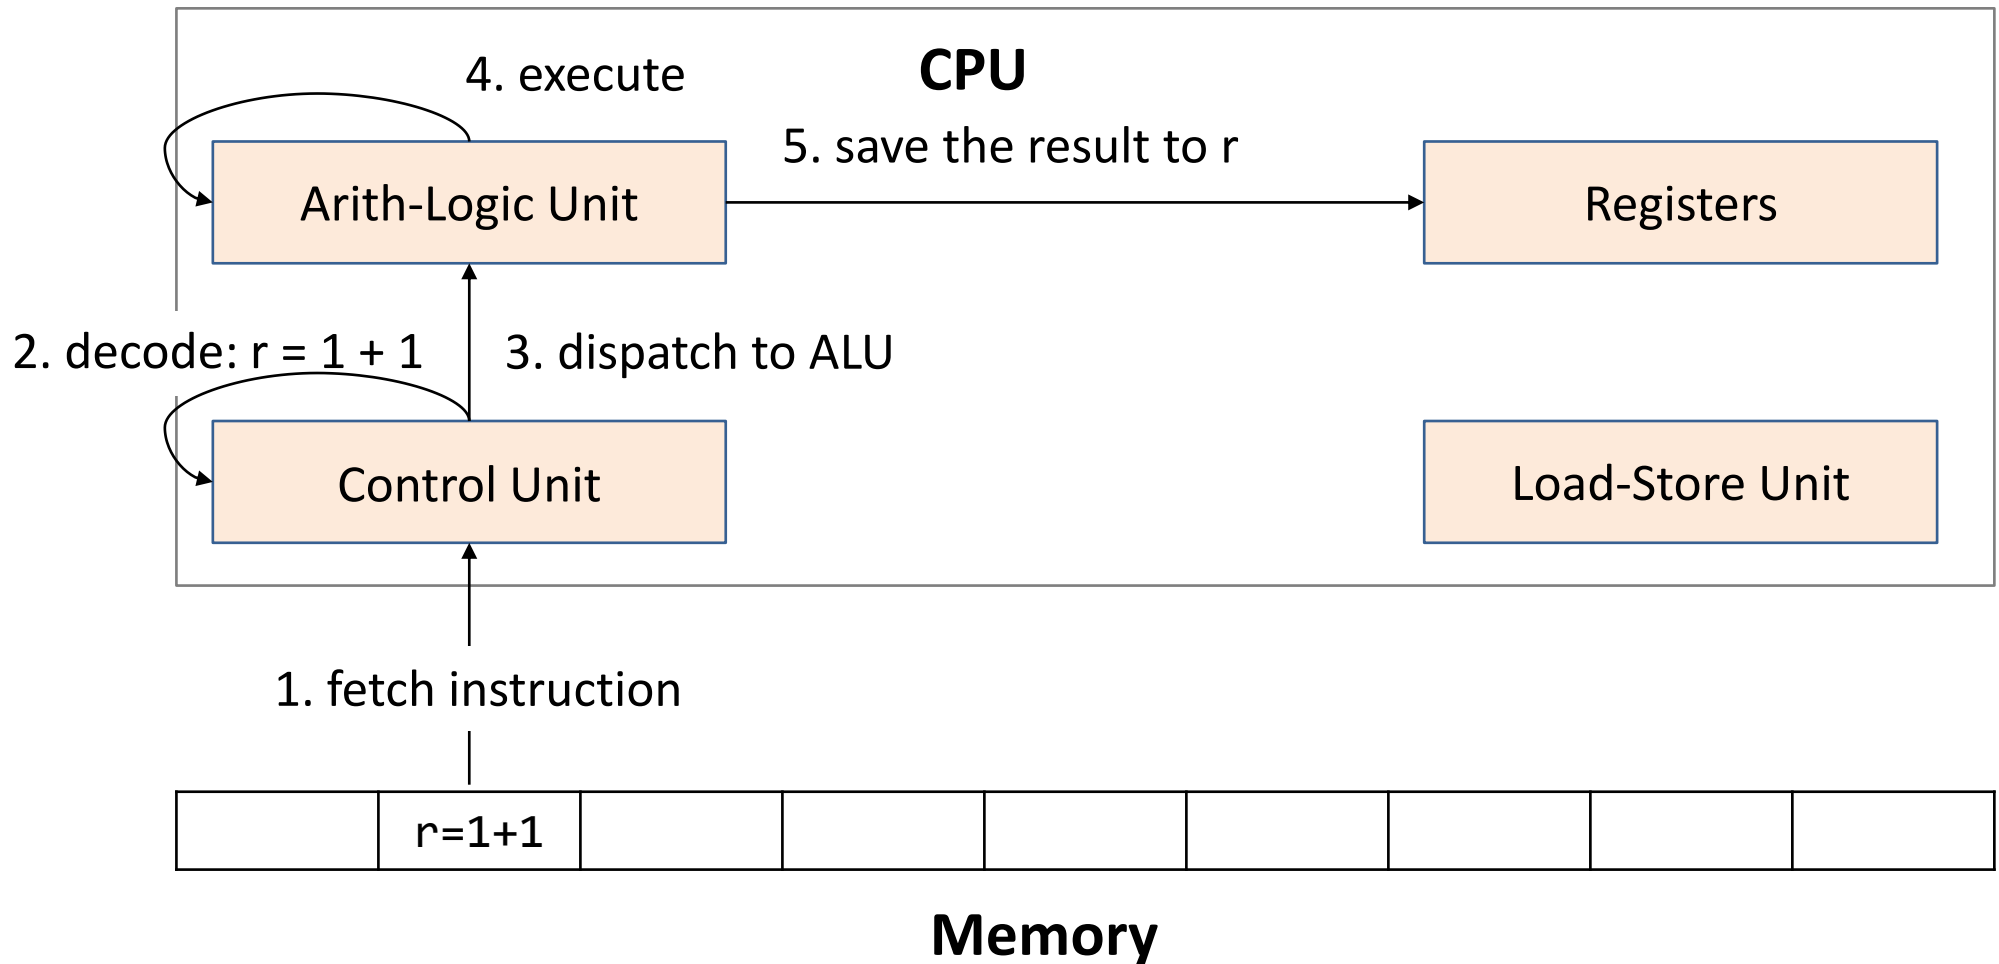
CPU: Central Processing Unit

中央处理器




- Arithmetic-logic Unit: perform arithmetic/logic operations
- Control unit: fetch instructions from memory; decode and dispatch
- Load-store unit: load/store data from/to memory
- Registers: supply operands to the ALU and save the execution result
- Memory: save data and instructions

Example: $r = 1 + 1$



Memory (Random-Access Memory or RAM)

- Read or write a data item based on its address in memory
- Each data unit generally takes 1 byte or 8 bits
- Access any data with the same latency irrespective of their locations



Address: 0x00 0x04 0x08 0x0B 0x10 0x14 0x18 0x1B 0x20

Example: $z = x + y$

```
z = x + y
```

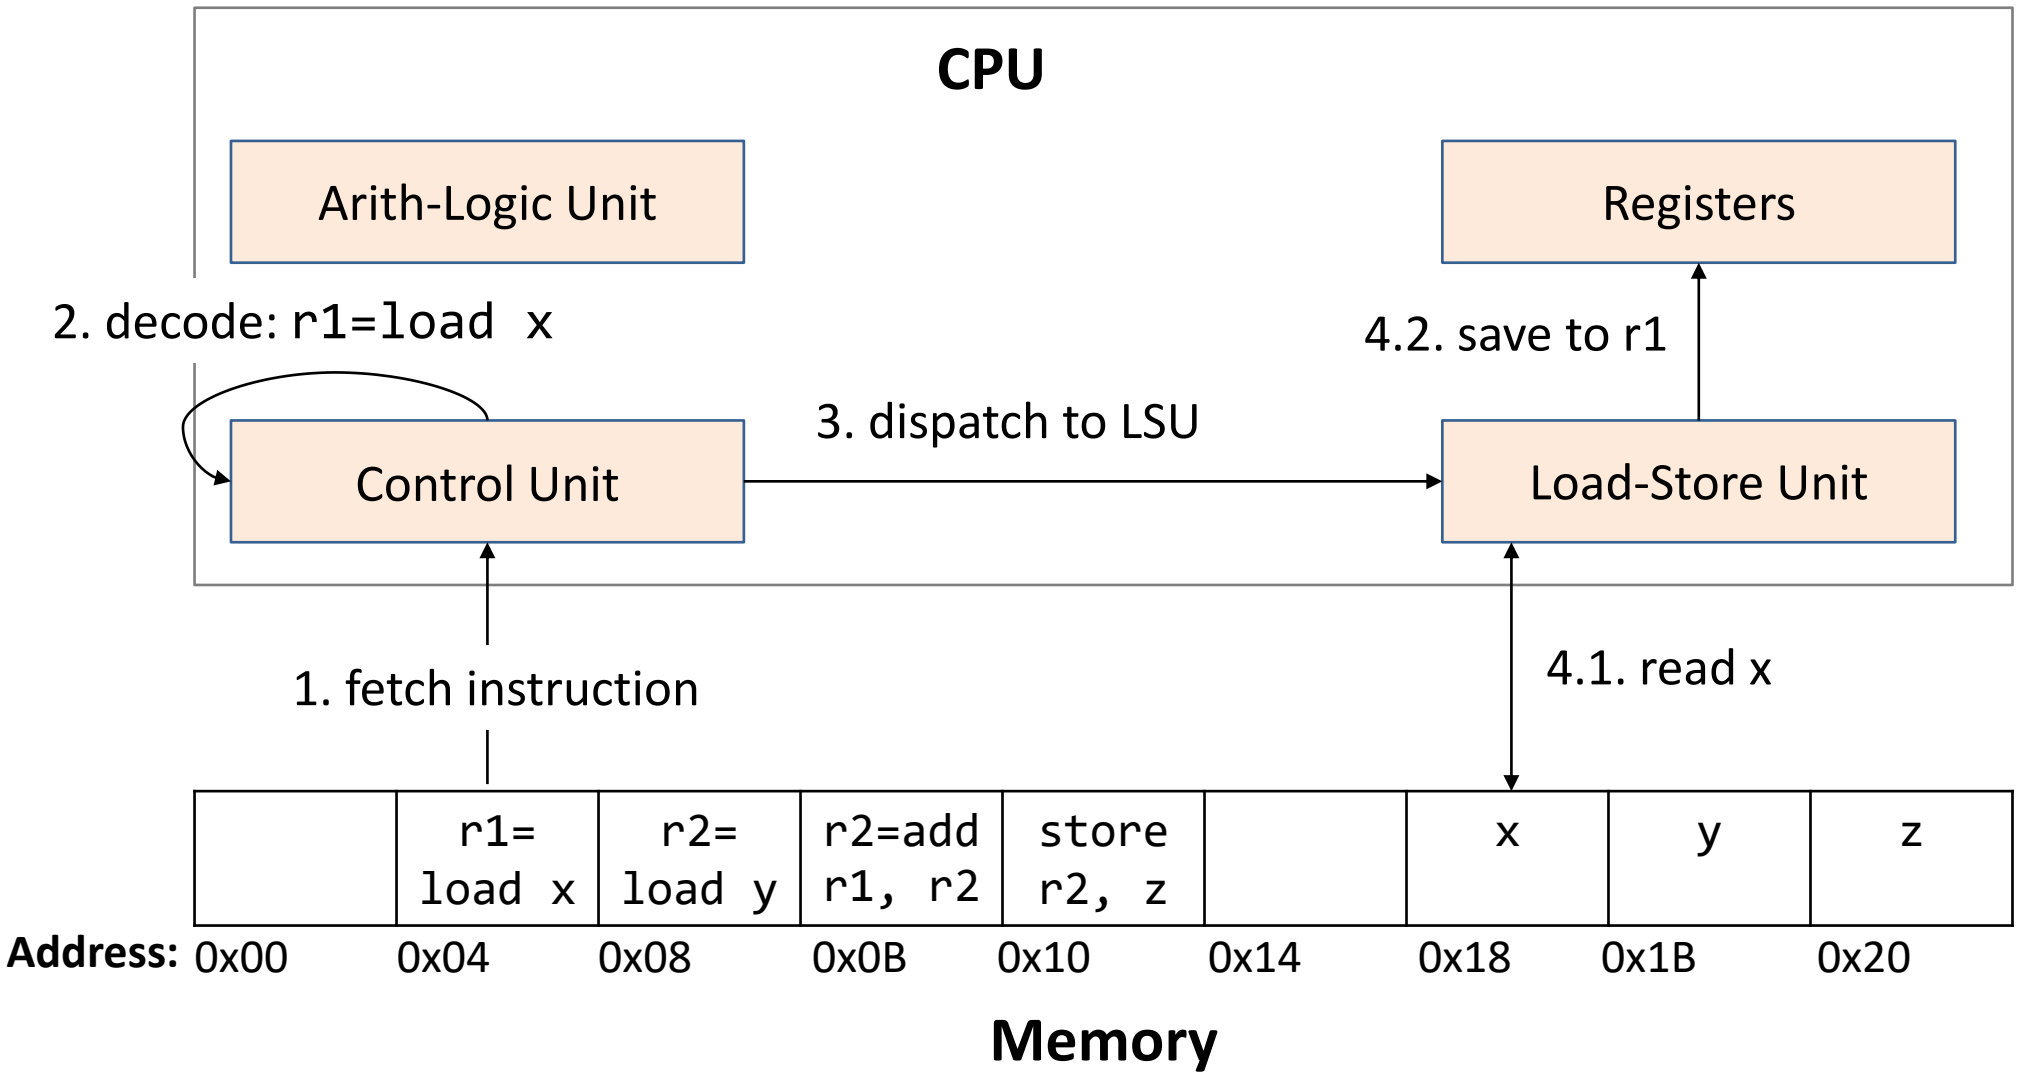
Compile
→
编译

```
r1 = load x  
r2 = load y  
r2 = add r1, r2  
store r2, z
```

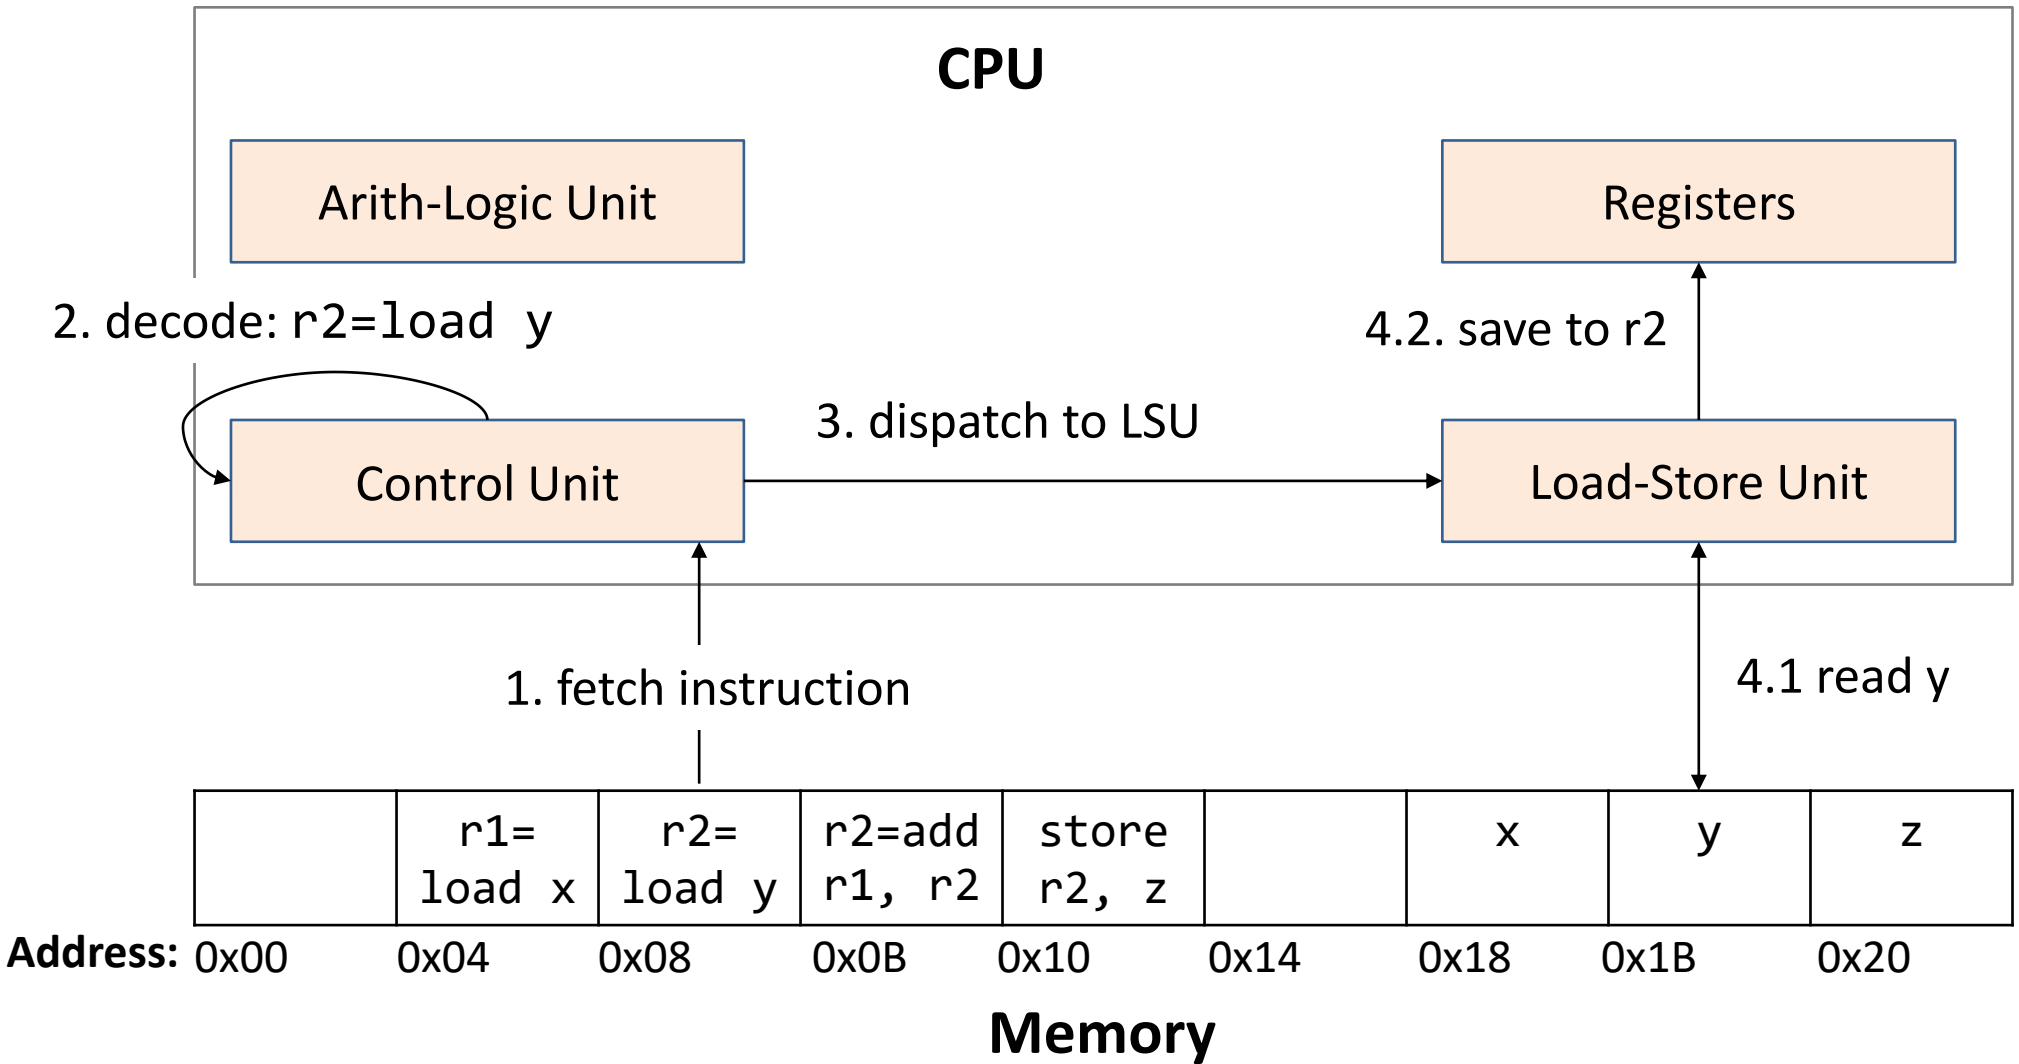
High-level Code
代码

Low-level Instructions
指令

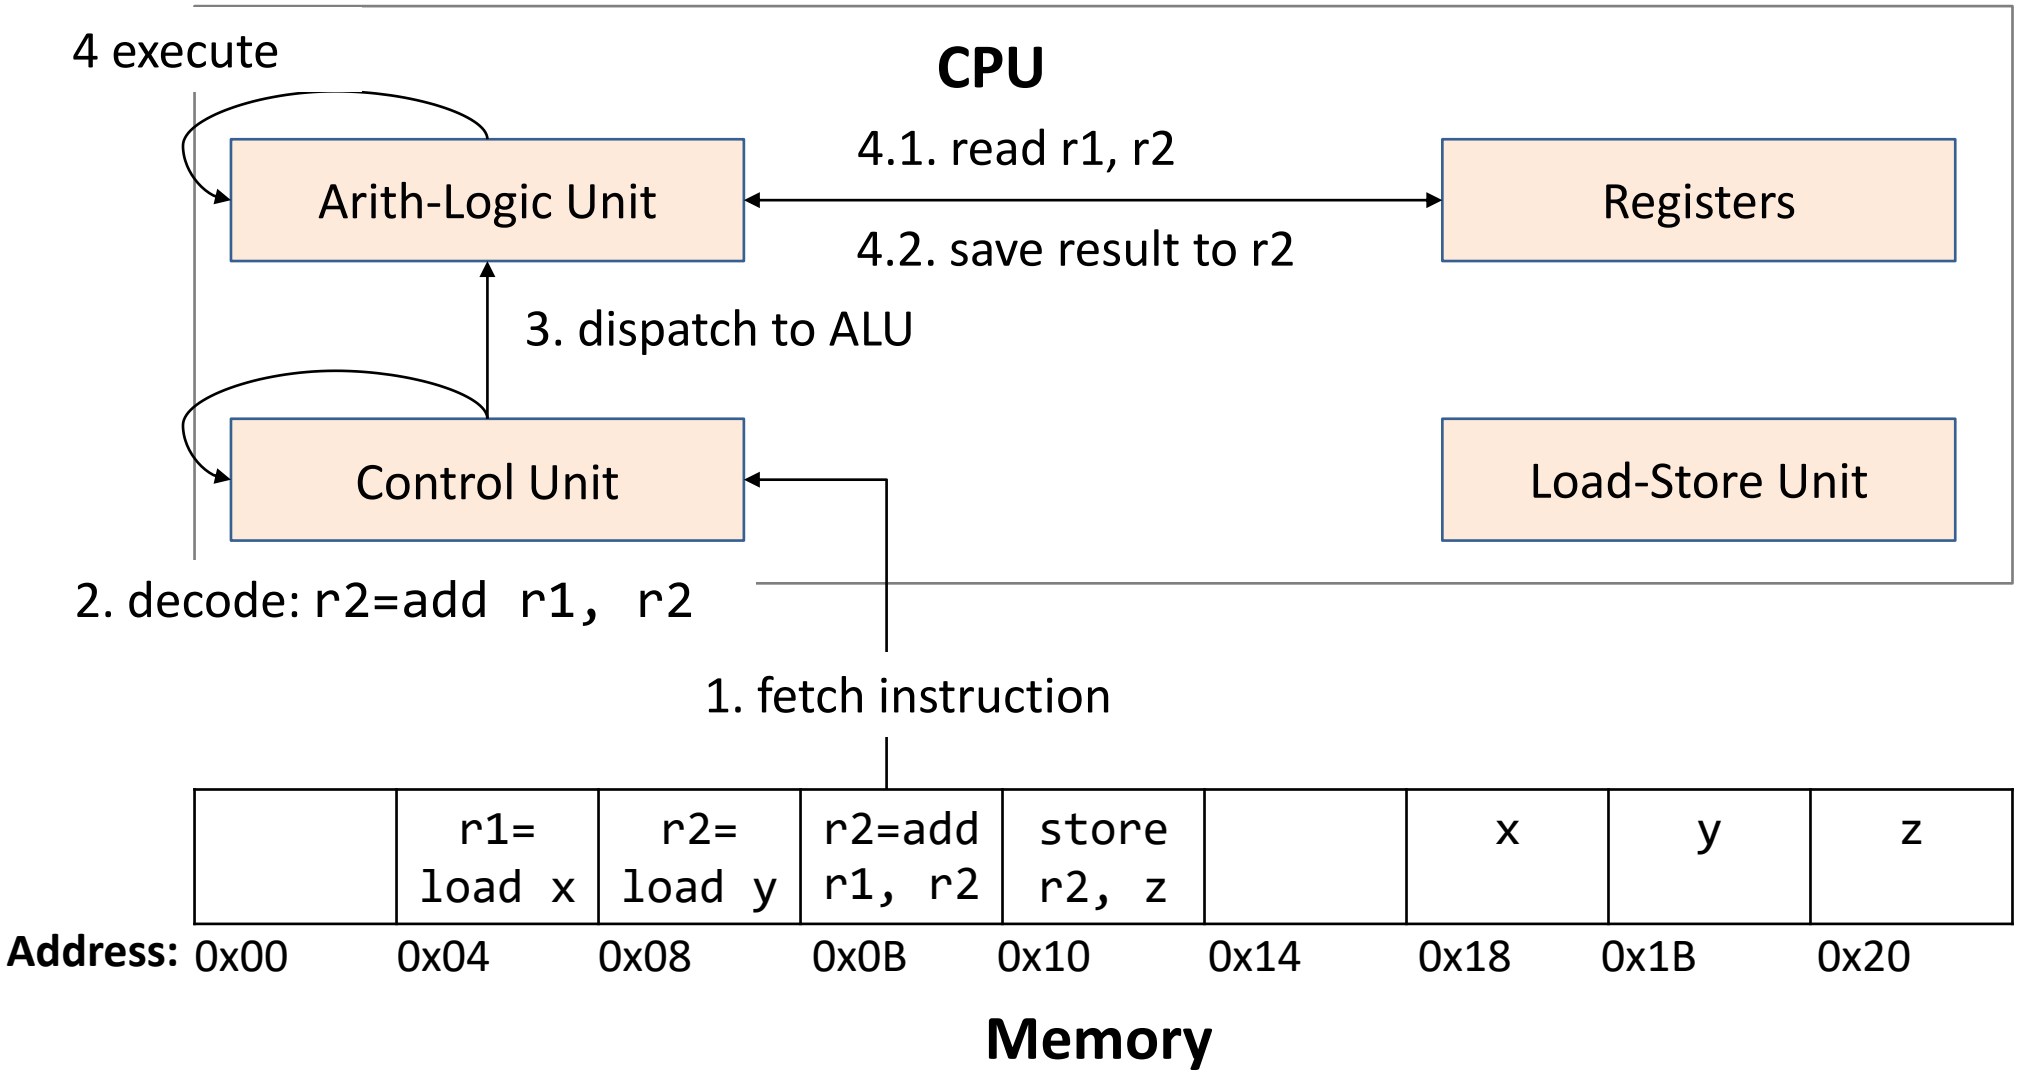
Step 1: r1 = load x



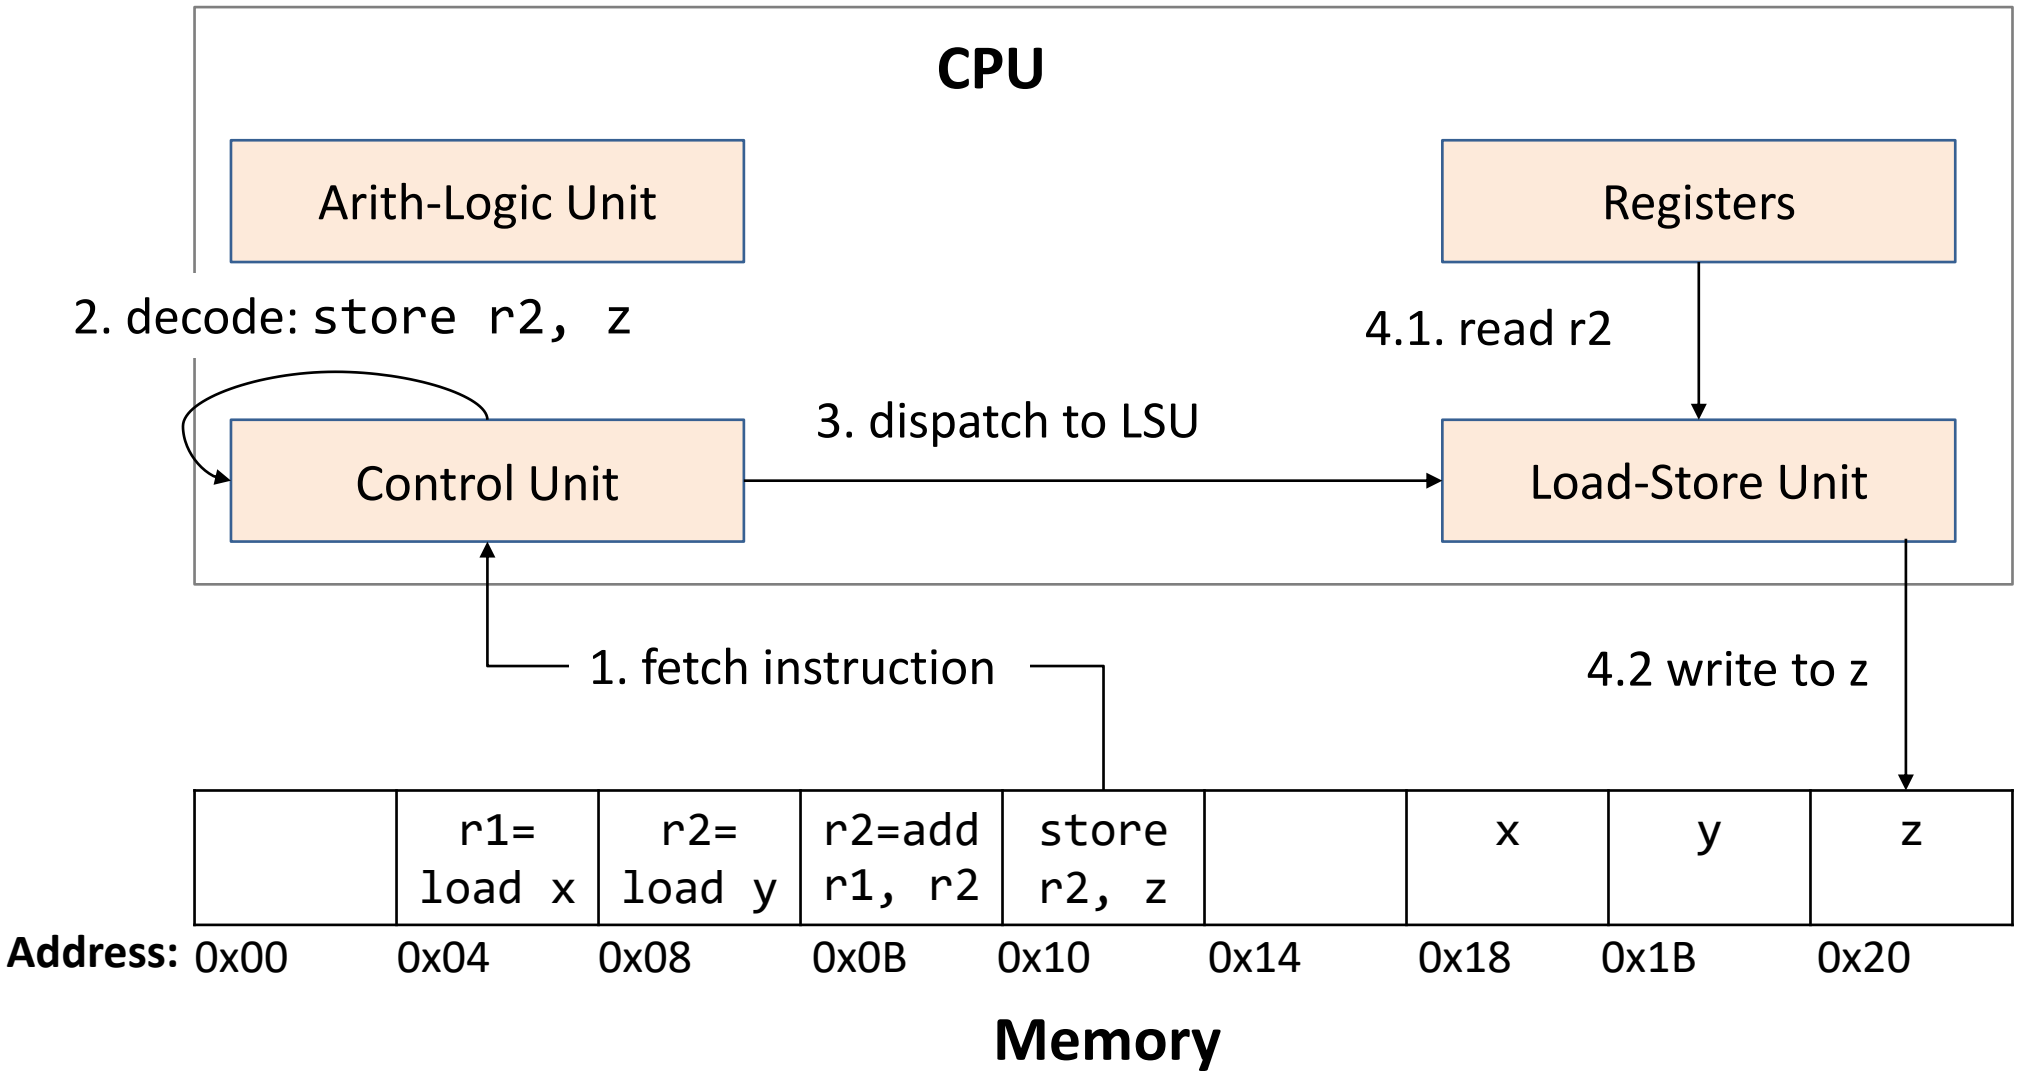
Step 2: r1 = load y



Step 3: r1 = load y



Step 4: store r2, z



Question

- If a CPU employs 32 bits to represent a memory address, what is the maximum memory space the CPU can support?

ISA: Instruction Set Architecture

指令集架构

- RISC: Reduced Instruction Set Computer
 - Separate data load/store and computation into different instructions
 - *e.g.*, AArch/ARM、RISC-V
- CISC: Complex Instruction Set Computer
 - An instruction can do both data load/store and computation
 - *e.g.*, X86, X86-64

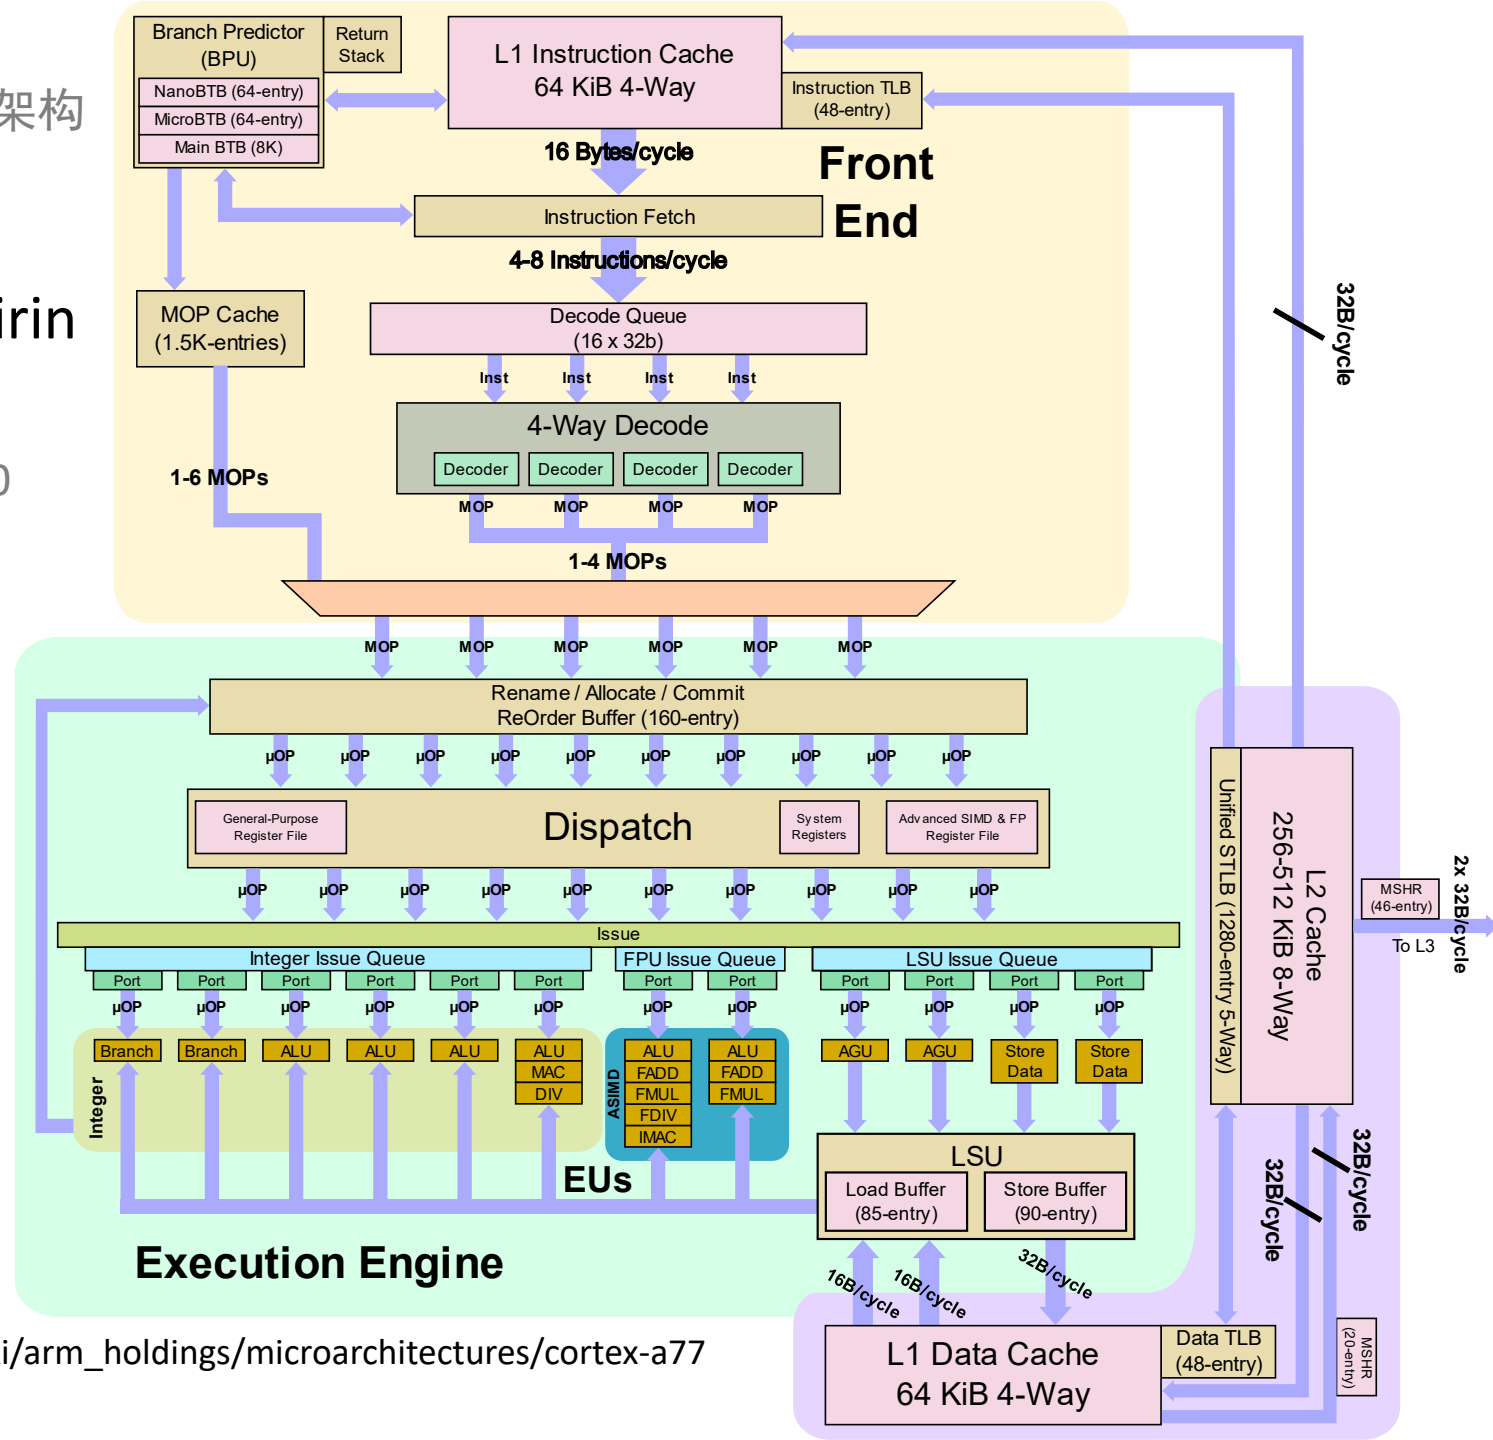
精简指令集

复杂指令集

Cortex-a77 微架构

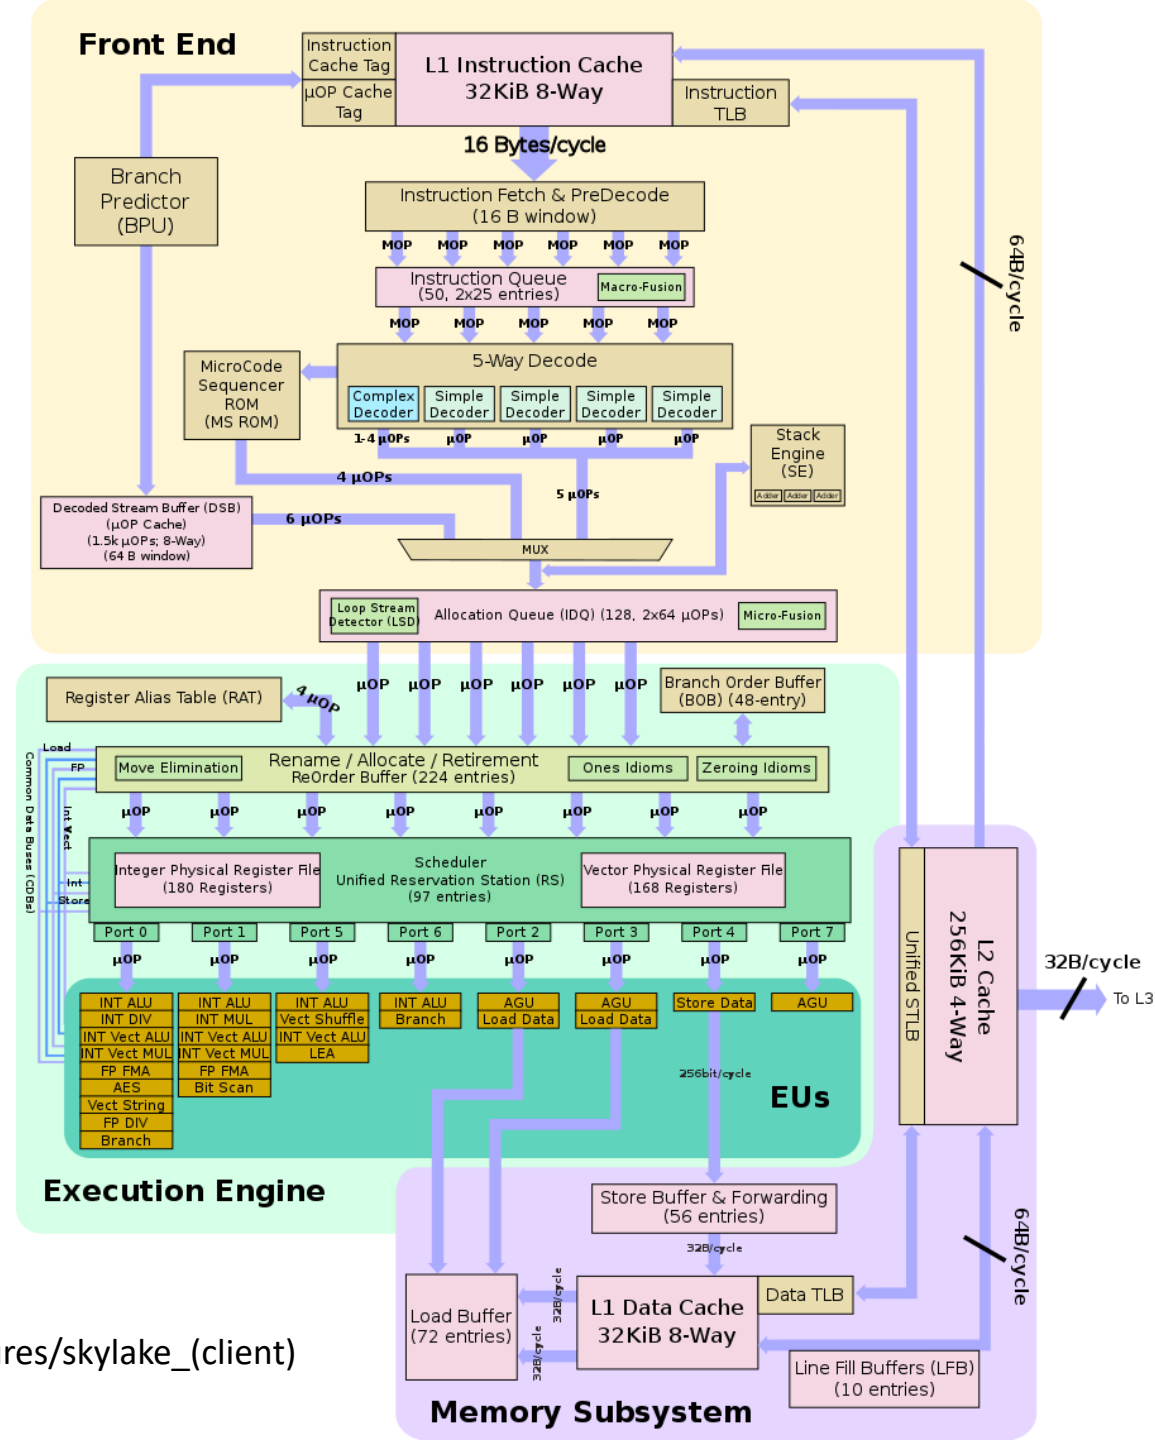
- ISA: ARMv8-A
- e.g., HiSilicon Kirin 9000

海思公司的麒麟9000



Skylake 微架构

- ISA: X86-64
- e.g., Intel i7, i9



ARMv8-A Instructions

```
mov x0, #5
mov x1, #10
add x2, x0, x1
sub x3, x0, x1
mul x4, x0, x1
sdiv x5, x0, x1
```

$x0 = 5$

$x1 = 10$

$x2 = x0 + x1$

$x3 = x0 - x1$

$x4 = x0 * x1$

$x5 = x0 / x1$

ARMv8-A Instructions: Data Load/Store

```
ldr x1, [x0]
```

$x1 = [x0]$

```
ldr x2, [x0, #12]
```

$x2 = [x0 + 12]$

```
add x3, x1, x2
```

$x3 = x1 + x2,$

```
str x3, [x0, #8]
```

$[x0 + 8] = x3$

```
add x2, x1, [x0]
```

Not allowed (CISC operation)

ARMv8-A Instructions: Control

控制指令

address Instructions

| | | |
|-------|---------|--------|
| 0x20 | cmp | x0, x1 |
| 0x24 | ble | .addr1 |
| ... | ... | |
| 0x104 | .addr1: | |
| ... | ... | |

[flag] = compare(x0, x1)



if(flag:<=) PC = 0x104

...

PC: 下一条指令地址


Instructions in Memory

- 32bit fixed-length for ARMv8-A 每条指令占固定内存空间
- Limitation: cannot employ some large numbers as operands



op S

add rn, rm, rd



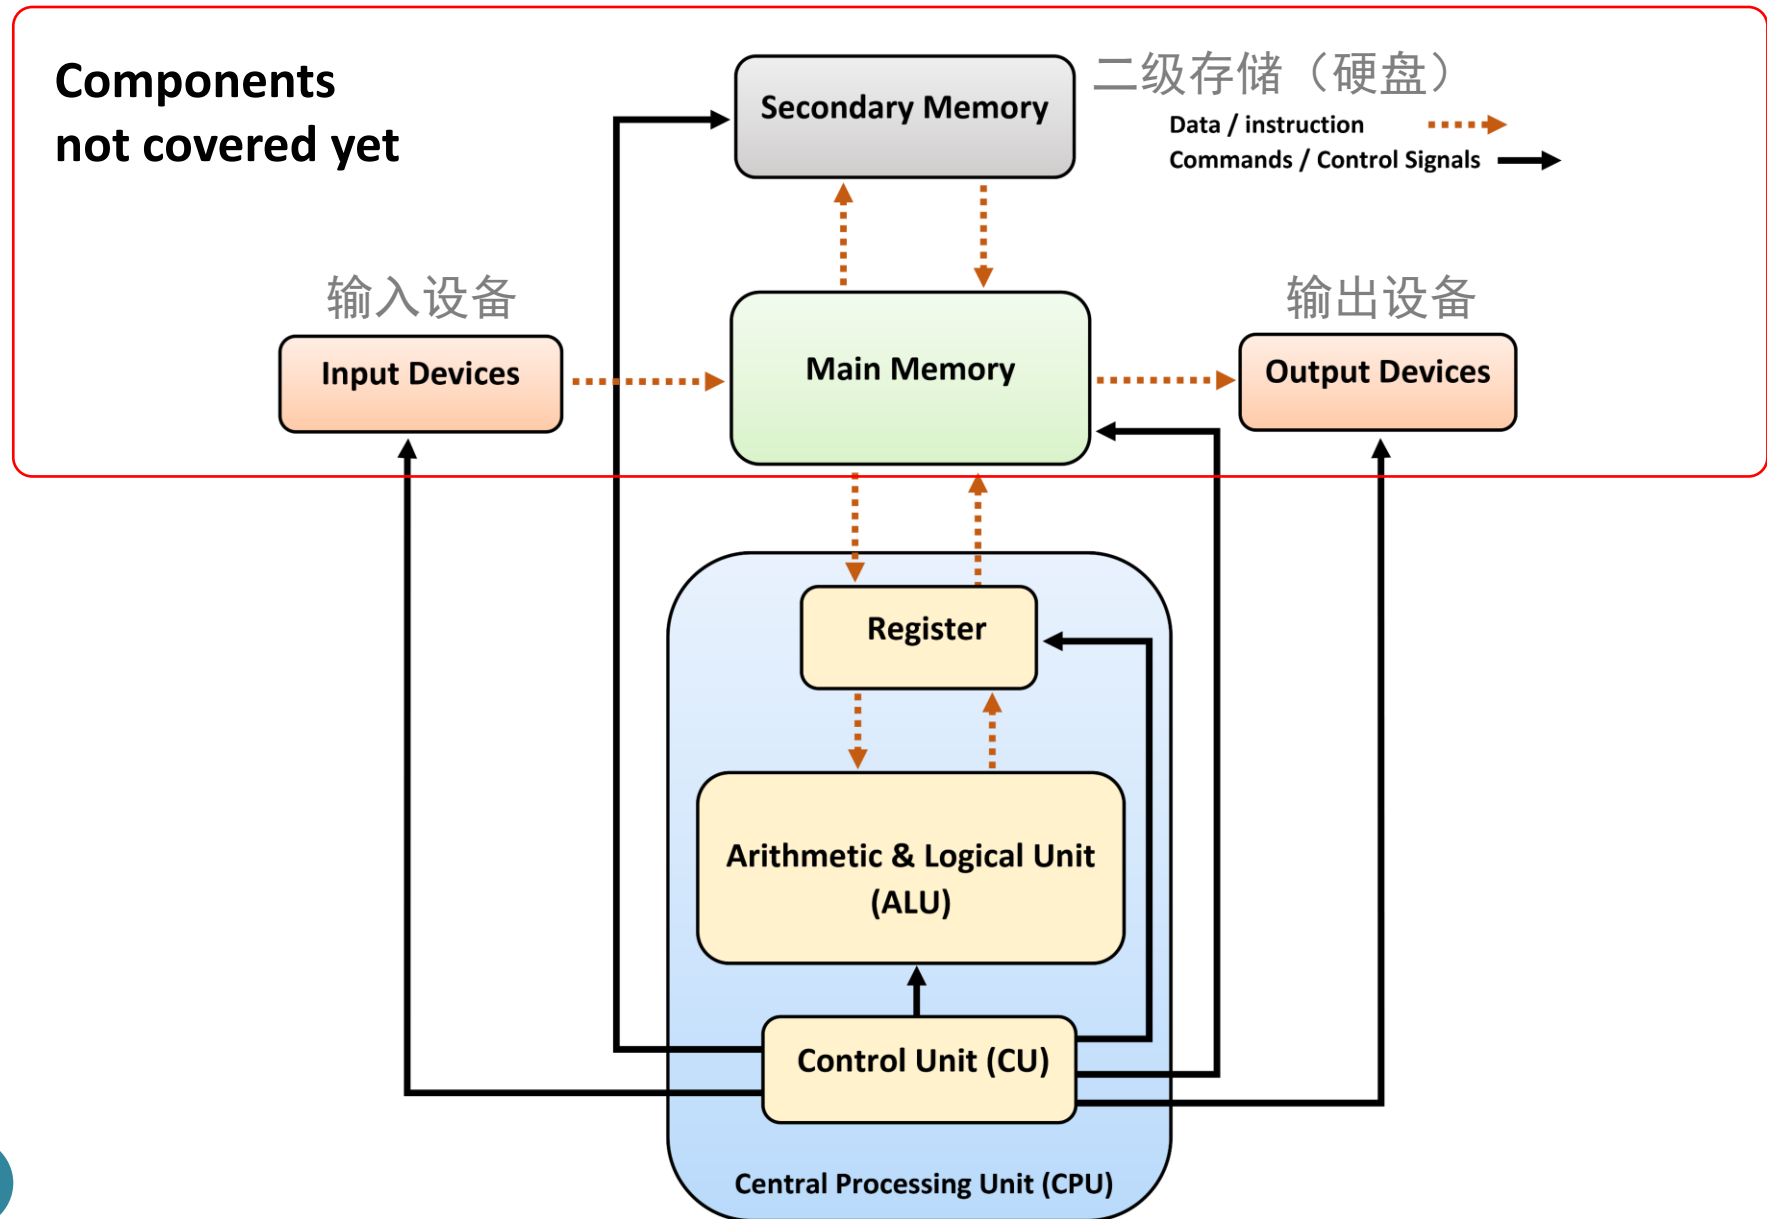
op S

add rd, rn, imm12
add rd, rn, #4095
add rd, rn, #4097

invalid

2. I/O and More

Computer Architecture



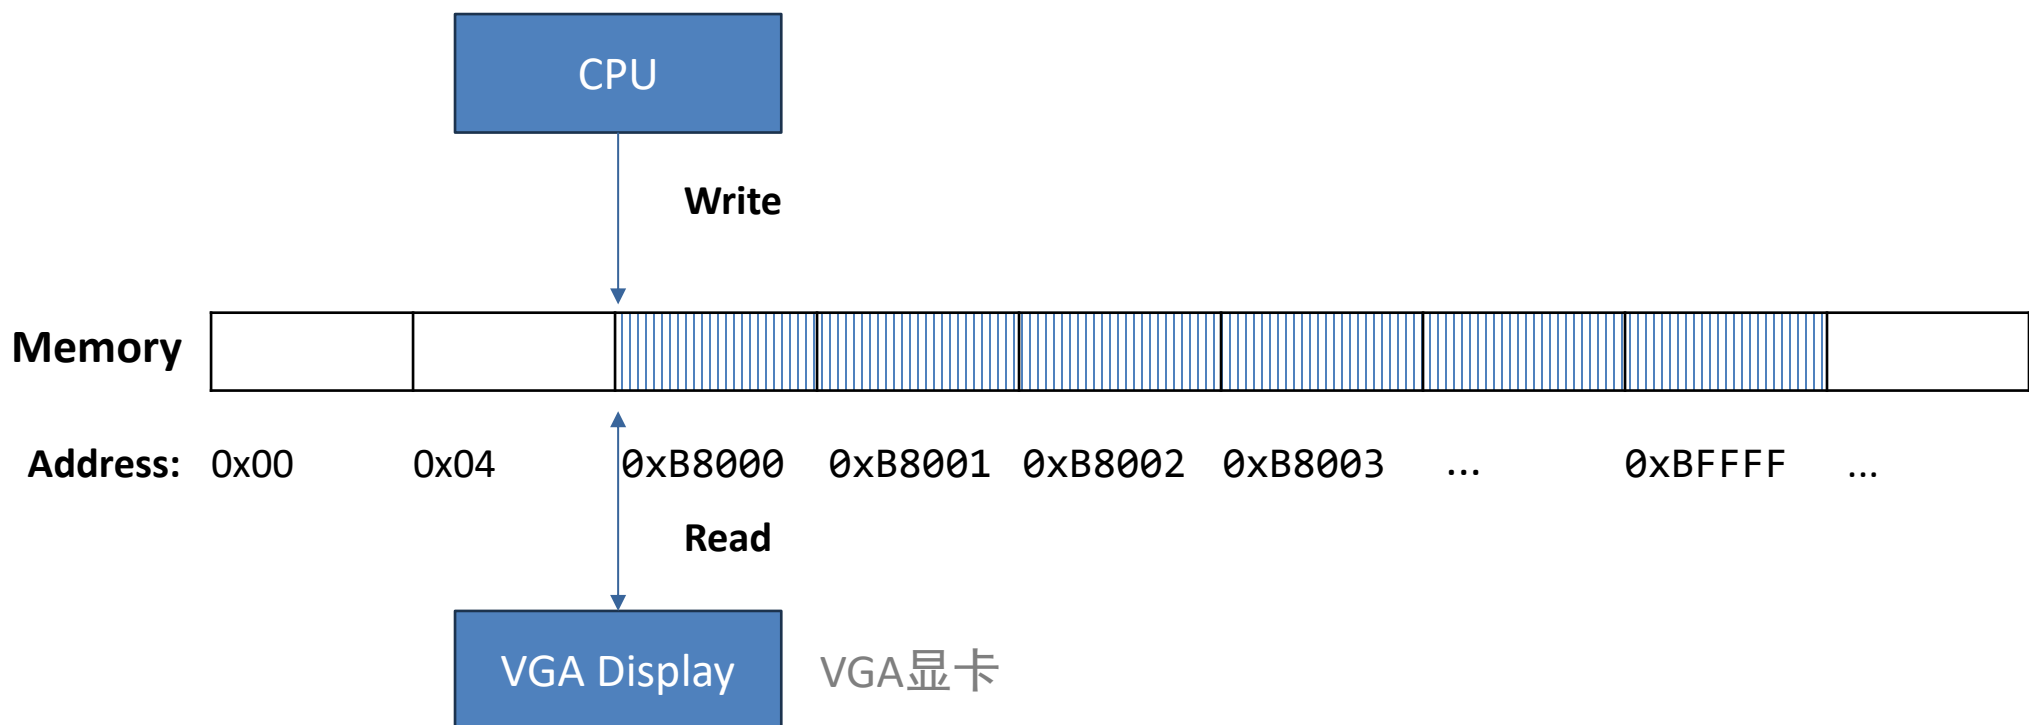
I/O Devices

- Input device: keyboard, mouse, *etc*
- Output device: display, printer, *etc*
- Question: how does a CPU handle I/O operations?

Option 1: Memory-mapped I/O

- Read/write data to I/O devices via memory access
- Addressing achieved via PCI (peripheral component interconnect)

外设互联总线



Output: Video as an Example

- VGA (Video Graphics Array): video display controller 15-pin
- Mapped to the PC's memory: 0xA0000 - 0xBFFFF:
- 0xB8000 for color text mode (32 KB)

```
mov ax, 0xB8000    ; 将显卡地址0xB8000保存到ax寄存器
mov es, ax         ; 将ax寄存器的值保存到es寄存器

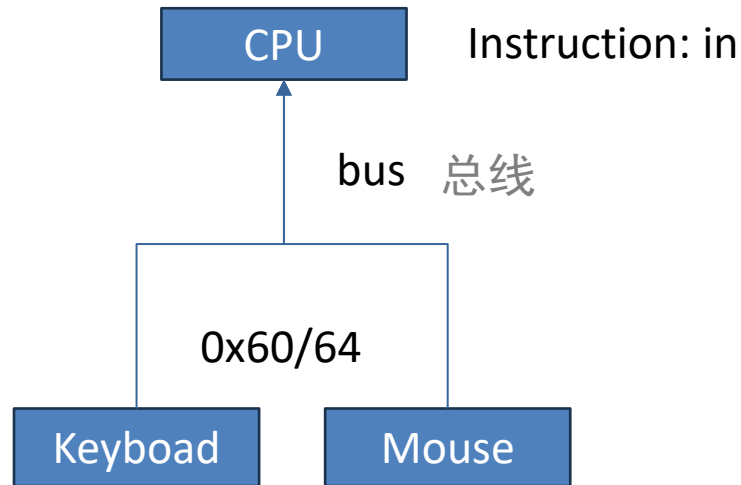
mov di, 0          ; 将显示位置0保存到di寄存器
mov al, 'A'        ; 将显示字符'A'保存到al寄存器
mov ah, 0x0F       ; 将显示颜色0x0F保存到ah寄存器
mov [es:di], al    ; 将al的值写如内存地址es+di
mov [es:di+1], ah  ; 将ah的值写如内存地址es+di+1
```

Question

- Why does the available physical memory not match the original purchase capacity?

Option 2: Port-based I/O

- Read/write data to I/O devices via port numbers.



Input is More Complicated

- Challenge: when/how many data arrives is not controlled by the CPU
 - Polling: let the CPU check particular signals periodically? How frequently?
 - By interrupt (trigger-based) via particular CPU pins

```
wait_for_input:
    in al, 0x64    ; 读取0x64端口信号到a1
    test al, 0x01 ; 测试a1是否为1
    jz wait_for_input ; 如果是0则继续轮询

read_input:
    in al, 0x60    ; 从0x60端口读入字符
```

Polling

轮询模式读取键盘输入

```
in al, 0x60
...
out 0x20, 0x20 ; 中断响应结束
```

Interrupt

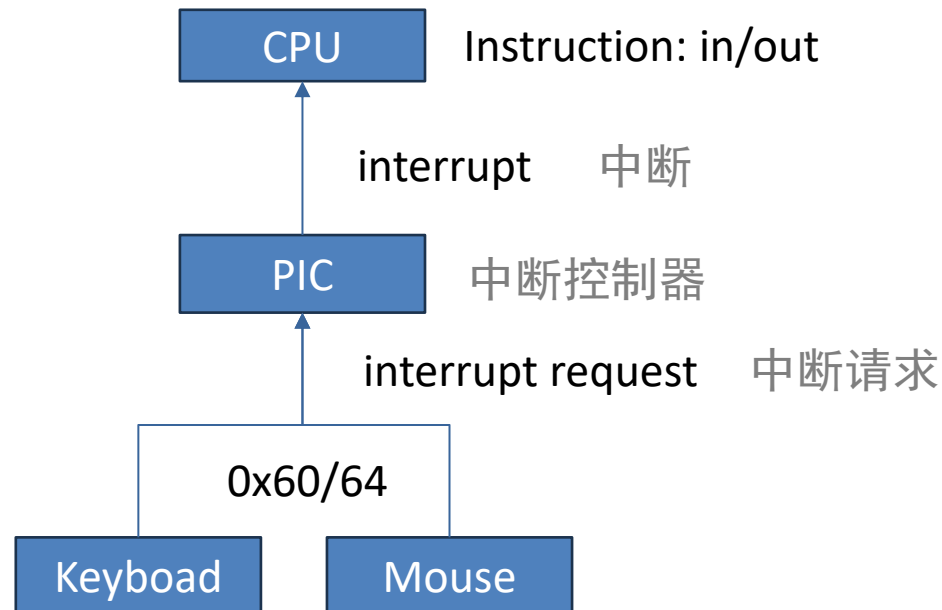
中断模式读取键盘输入



Input is More Complicated

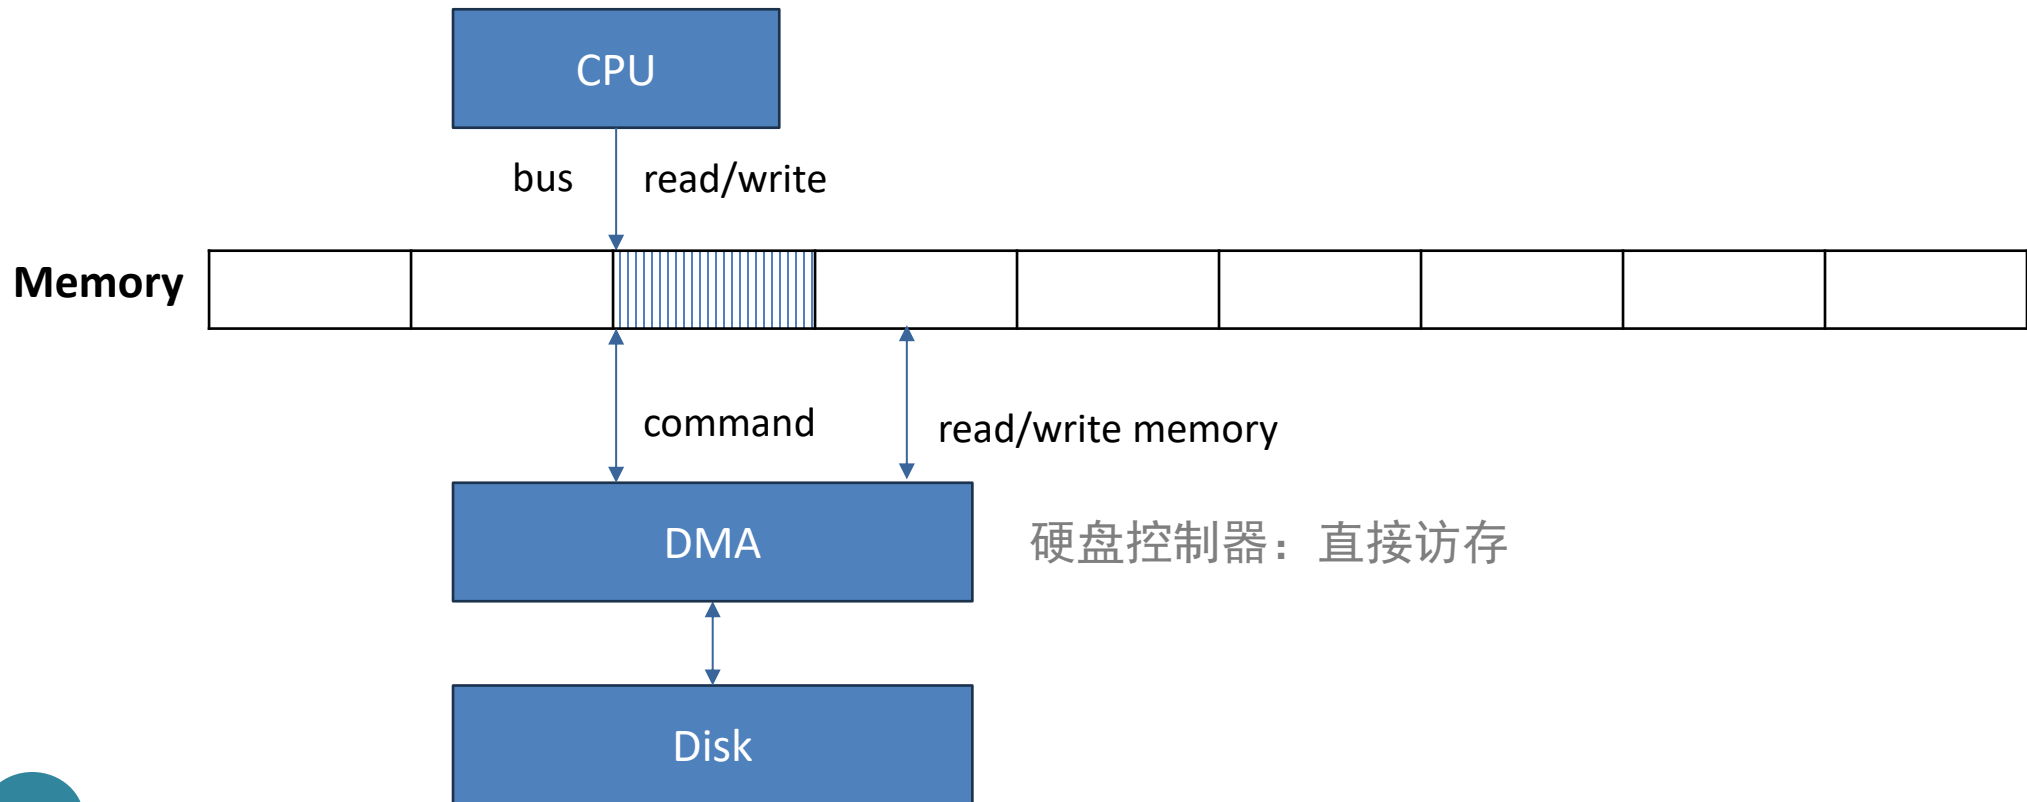
- What if data arrives from multiple input devices simultaneously?
 - First come first serve? May lead to denial-of-service
 - Priority/queuing via PIC (programmable interrupt controller)

优先级 排队



Disk I/O

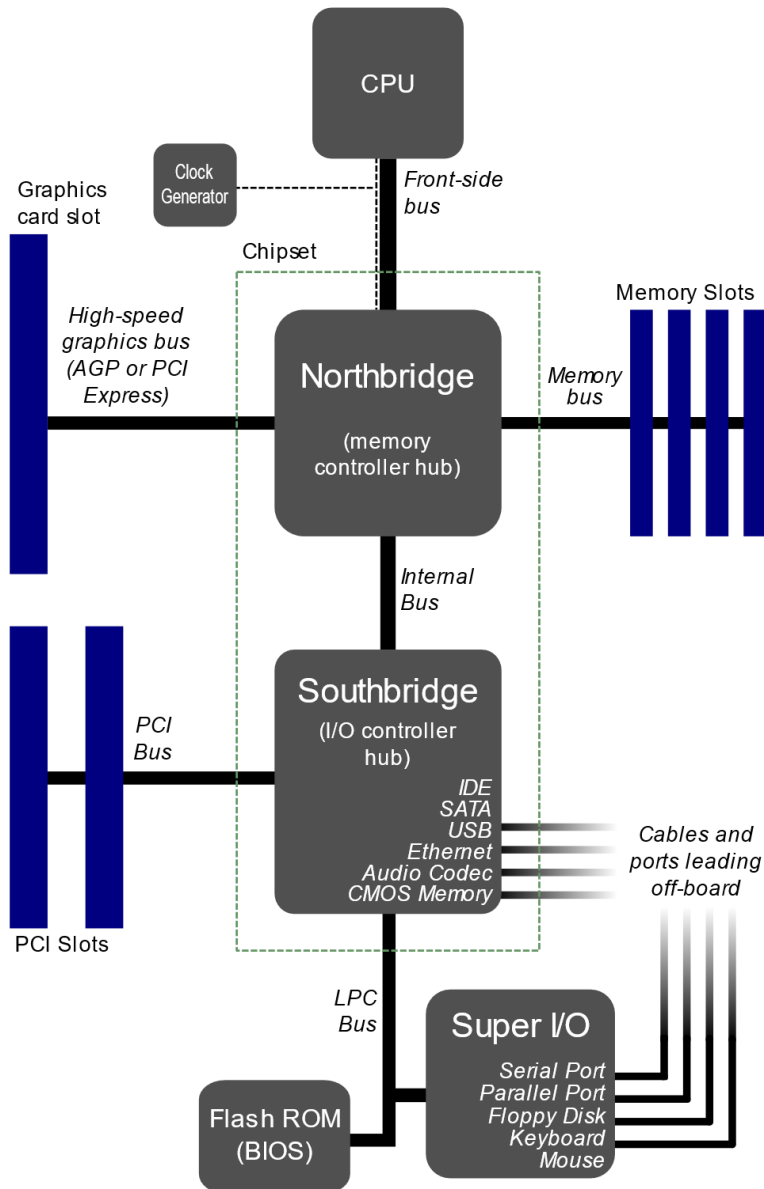
- Disk read/write operations are much more slower than memory
- Offload the workload to disk controller via direct memory access



Question


- Why does the available disk space not match the capacity at the time of purchase?

Motherboard 主板



- Northbridge: handle high-performance tasks, connected directly to a CPU.
- Southbridge: handle low speed tasks, connected via the northbridge

Motherboard 主板



Several Concepts

- Hardware: physical components of a computer system
- Software: programs that run on the hardware
 - *e.g.*, operating systems and applications
- Firmware: specialized software embedded into hardware devices to control or manage the hardware directly
 - *e.g.*, BIOS (basic input/output system) for interrupt handling

3. In-class Practice

In-class Practice 1

- What is the CPU info and RAM size of your PC?
- Show the screenshot and highlight the corresponding region



Windows: Win + R: msinfo32.exe



MacOS

In-class Practice 2

- The following code computes the largest Fibonacci number less than 100 and saves the result to the register x0 (aarch64) or eax (x86).
- Modify the X86 version to compute the largest accumulated value less than 100, starting from 1.

```
_main:
    mov x0, #1
    mov x2, #1

_loop:
    mov x1, x0
    mov x0, x2
    add x2, x2, x1
    cmp x2, #100
    ble _loop
    ret
```

Aarch64 Version



```
_main:
    mov eax, 1
    mov ecx, 1




_loop:
    mov ebx, eax
    mov eax, ecx
    add ecx, ebx
    cmp ecx, 100
    jle _loop
    ret
```

X86 Version


ecx = ecx + ebx

In-class Practice 2

 **Assembly x86 Emulator**

 Run  Step  Reset

```
1  _main:
2      mov  eax, 1
3      mov  ecx, 1
4
5  _loop:
6      mov  ebx, eax
7      mov  eax, ecx
8      add  ecx, ebx
9      cmp  ecx, 100
10     jle  _loop
11     ret
12
```

Registers 

| | | | |
|-----|------------|--------------|------------|
| EAX | 0x00000059 | EBX | 0x00000037 |
| ECX | 0x00000090 | EDX | 0x00000000 |
| ESI | 0x00000000 | EDI | 0x00000000 |
| EBP | 0x00000000 | ESP | 0x00020400 |
| EIP | 0x0002041C | Ln 11, Col 5 | |

Flags

| | | | |
|-------|---|----------|---|
| Carry | 0 | Dir | 0 |
| Int | 0 | Overflow | 0 |
| Sign | 0 | Zero | 0 |