

MF20006: Introduction to Computer Science

Lecture 5: Data Structures

Hui Xu

xuh@fudan.edu.cn



Outline

1. Basic Concept

2. Linear Data Structures

3. Trees and Graphs

4. Hash Set and Hash Map

1. Basic Concept

How to Describe Companies with a Program?

HOME > NVDA · NASDAQ

NVIDIA Corp

\$116.00 ↓ 10.33% -13.37 1M

Sep 23, 12:31:01 AM UTC-4 · USD · NASDAQ · Disclaimer

1D 5D 1M 6M YTD 1Y 5Y MAX



[Compare to](#)

Tesla Inc

\$238.25

TSLA ↑ 8.14%

Apple Inc

\$228.20

AAPL ↑ 0.60%

Microsoft Corp

\$435.27

MSFT ↑ 4.43%

Amazon.com Inc

\$191.60

AMZN ↑ 8.22%

PREVIOUS CLOSE

\$117.87

DAY RANGE

\$115.39 - \$118.62

YEAR RANGE

\$39.23 - \$140.76

MARKET CAP

2.85T USD

AVG VOLUME

332.34M

P/E RATIO

54.48

DIVIDEND YIELD

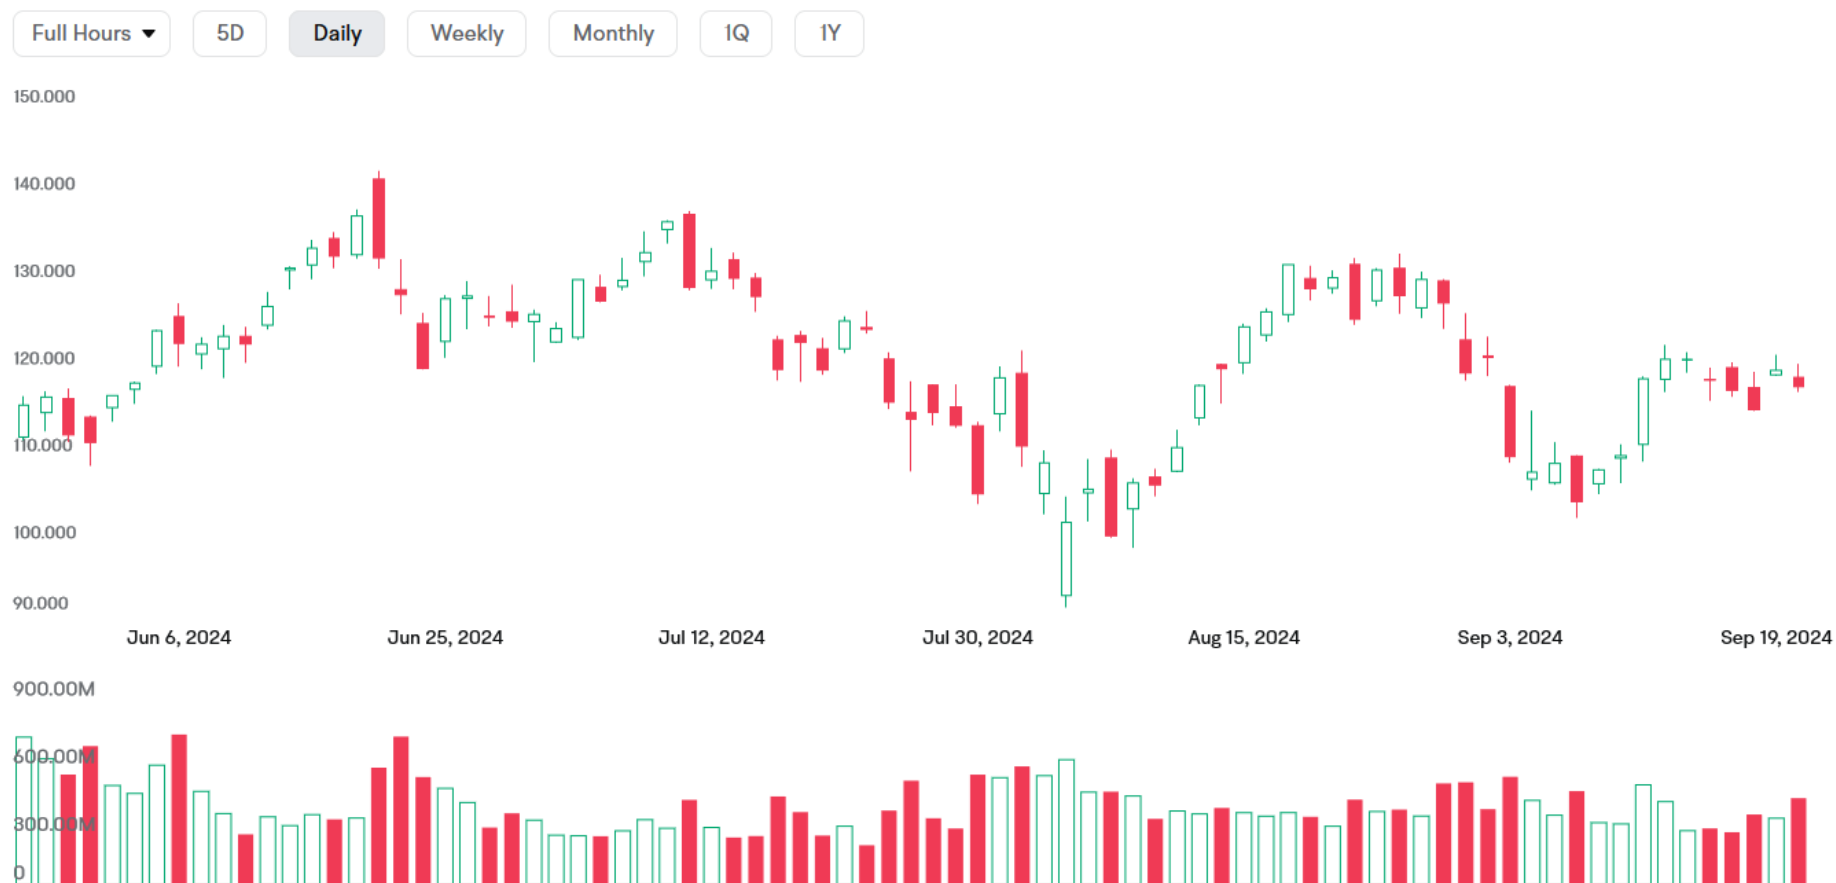
0.03%

PRIMARY EXCHANGE

NASDAQ

How to Manage Historical Prices?

❑ We want to collect the historical price, *e.g.*, as a candlestick chart



Data Structure

- ❑ **A specialized format for organizing and storing data, with defined operations for accessing and modifying it efficiently.**
 - How data is arranged in memory.
 - How we insert, delete, search, and update elements.

Implement Your Own Data Structure with Dataclass

❑ A dataclass is a special kind of Python class that automatically gives you the usual methods you'd write for a data-holding class.

❑ It generates an `__init__()` method automatically.

➤ The field type should be specified.

```
from dataclasses import dataclass

@dataclass
class Point():
    x: int
    y: int

    def other_methods():
        ...
```



```
class Point():

    def __init__(self, x, y):
        self.x = x
        self.y = y

    def other_methods():
        ...
```

If the field type is not specified...

- ❑ The class constructor does accept the field parameter.

```
@dataclass
class Point():
    x: int
    y: int

let p = Point(10, 20)
```

=

```
@dataclass
class Point():
    x = 0
    y = 0

let p = Point()
p.x = 10
p.y = 20
```


Data Structure of a Company

```
@dataclass
class Company:
    name: str          # "NVIDIA Corporation"
    ticker: str         # "NVDA"
    exchange: str       # "NASDAQ"

    # Trading info
    current_price: float # $100.1
    open_price: float    # $99.12
    close_price: float
    high_price: float
    low_price: float     # $98.79
    volume: int          # 1000

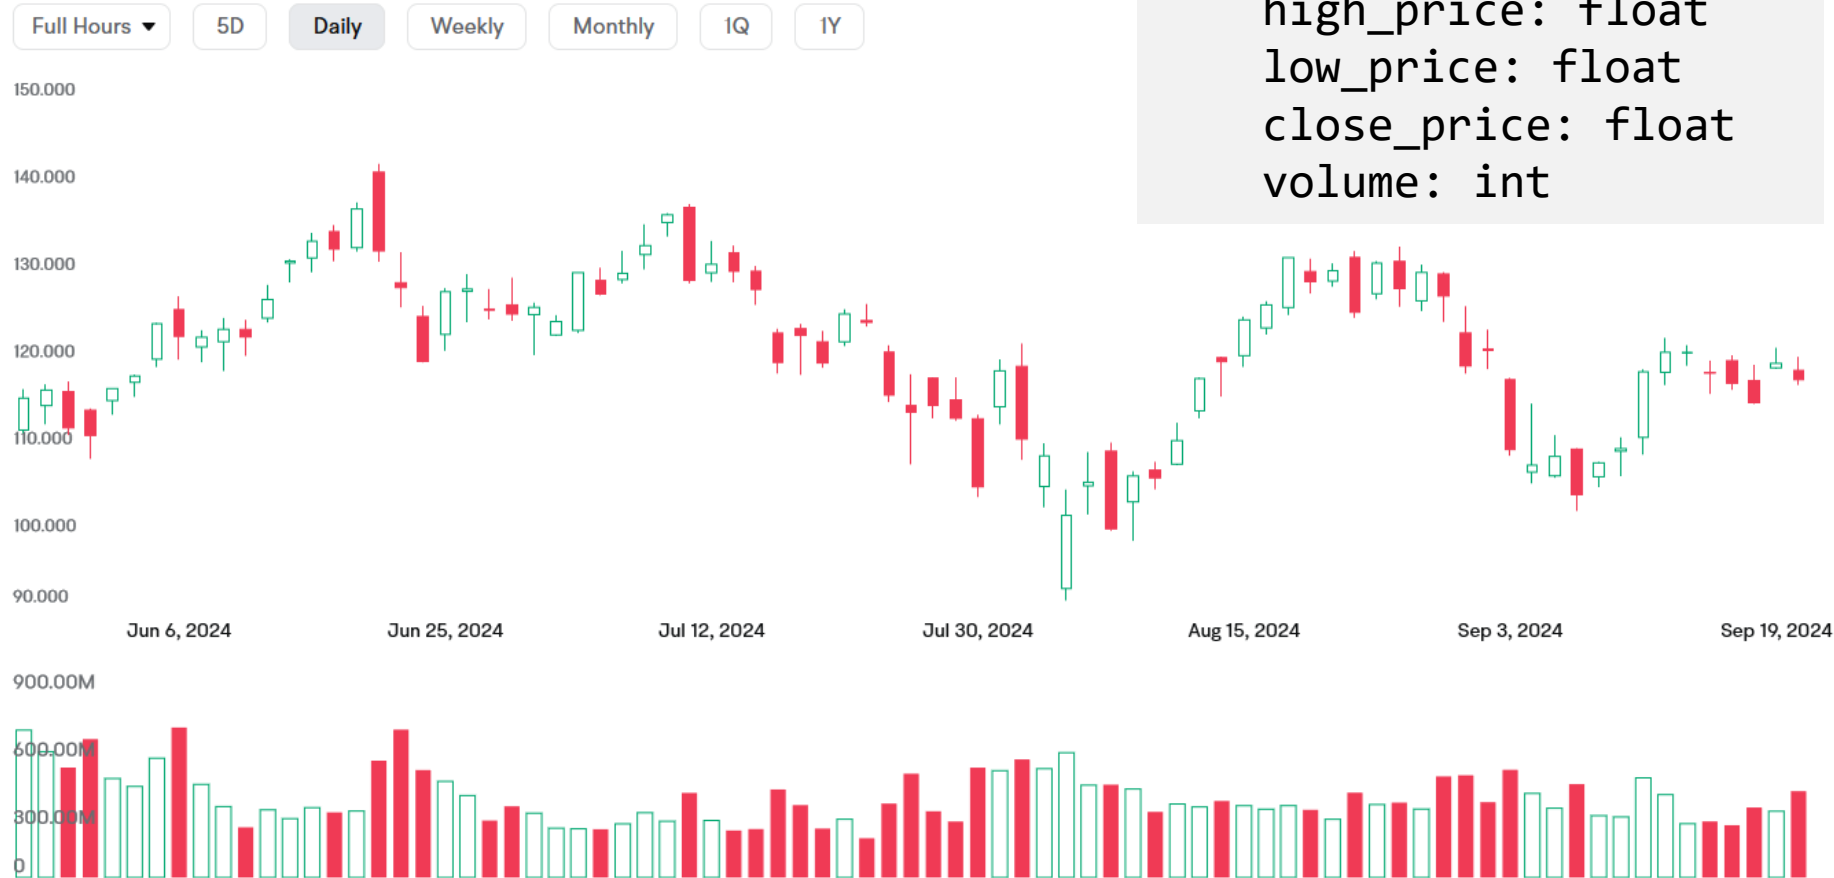
    # Financial info
    pe: int              # P/E ratio
    pb: float            # P/B ratio
    market_cap: float   # e.g., $1000 billion
```

Data Structure of Historical Prices

❑ We want to collect the historical price, *e.g.*, as a candlestick chart

```
@dataclass
class Company:
    ...
    historical_prices: List[TradingInfo]
```

```
@dataclass
class TradingInfo:
    date: int
    open_price: float
    high_price: float
    low_price: float
    close_price: float
    volume: int
```



2. Linear Data Structures

Linear Data Structures

❑ Consider the distribution of data:

- Array: a collection of elements stored in contiguous memory locations.
- Linked List: data are not contiguously stored.

❑ Consider the behaviors:

- Queue: first-in-first-out.
- Stack: first-in-last-out.

Array

- ❑ A collection of elements stored in contiguous memory locations.
- ❑ All elements are of the same data type.
- ❑ Length: number of elements within the array.
- ❑ Size: means the memory space it occupied.

Memory Address	0x200	0x201	0x202	0x203	0x204	0x205	0x206	0x207	0x208	0x209
Data	F	I	S	F	1	3	0	0	2	0
Index	0	1	2	3	4	5	6	7	8	9

Array

- ❑ Access any elements via base address + offset.
- ❑ Supposing the size of each data unit is 4 types, *e.g.*, 32bit integer, we can retrieve the *i*th data from the memory address via `a[i]`.

$$a[i] = a[0] + 4*i$$

Memory Address	0x200	0x204	0x208	0x20c	0x210	0x214	0x218	0x21c	0x220	0x224
Data	1	1	2	3	5	8	13	21	34	55
Index	0	1	2	3	4	5	6	7	8	9

Array length: 10

Array size: 40 bytes

Array Operations

❑ **Read/write any elements via base address + offset.**

➤ Constant time, denoted as $O(1)$.

❑ **Searching an element from an array of length n .**

➤ Average: $n/2$ elements until find it, denoted as $O(n)$.

❑ **Insertion or deletion an element requires shifting the rest elements.**

➤ Average: moving $n/2$ elements, denoted as $O(n)$.

Memory Address	0x200	0x204	0x208	0x20c	0x210	0x214	0x218	0x21c	0x220	0x224
Data	1	1	2	3	5	8	13	21	34	55
Index	0	1	2	3	4	5	6	7	8	9

Two-dimensional Array: Matrix

- ❑ Consist of multiple one-dimensional array; each has the same length.
- ❑ Supposing an i32 array has m rows, and each row has length n, we can retrieve data of the *i*th row and *j*th column via `a[i][j]`.

$$a[i][j] = a[0][0] + 4*i*n + j$$

Row Index	0	1	1	2	3	5
	1	8	13	21	34	55
		0	1	2	3	4
		Column Index				


Array View

Memory Address	0x200	0x204	0x208	0x20c	0x210	0x214	0x218	0x21c	0x220	0x224
Data	1	1	2	3	5	8	13	21	34	55
Index	0, 0	0, 1	0, 2	0, 3	0, 4	1, 1	1, 2	1, 3	1, 4	1, 5


Memory View

Linked List

- ❑ Similar as array, but the data are not contiguously stored.
- ❑ Each list node has a data field and an address field to the next or the previous node.



Single-linked List



Double-linked List

Pros and Cons of List Operations

- ❑ **The cost of accessing elements of different positions varies a lot.**
 - Related to the length of the list: denoted as $O(n)$.
- ❑ **Insertion/deletion an element at any position costs constant time.**
 - Unrelated to the length of the list: $O(1)$.

Does Python have Array or Linked List?

❑ **Python's built-in list is neither linked list nor array.**

- The built-in list is a dynamic array.
- Each array element is an object pointer.
- Elements can be different types.

❑ **Alternative feature: `numpy.ndarray`.**

Example: Matrix Addition with Python

```
def add(a: list[list[int]], b: list[list[int]]) -> list[list[int]]:
    rows = len(a)
    cols = len(a[0])
    result = [[0 for _ in range(cols)] for _ in range(rows)]
    for i in range(rows):
        for j in range(cols):
            result[i][j] = a[i][j] + b[i][j]
    return result
```

Matrix Addition via NumPy Native Code

- ❑ Python executes each statement through the interpreter.
- ❑ In comparison, NumPy native code bypasses the interpreter.
 - Compiled directly into machine instructions.

```
# Create two 2x2 arrays
a = np.array([[1, 2],
              [3, 4]])
b = np.array([[5, 6],
              [7, 8]])

c = a + b
```

Benchmark the Performance

```
size = 1000 # 1000x1000 matrix
a_list = [[random.random() for _ in range(size)] for _ in range(size)]
b_list = [[random.random() for _ in range(size)] for _ in range(size)]
a_np = np.array(a_list)
b_np = np.array(b_list)

start = time.perf_counter()
add(a_list, b_list)
end = time.perf_counter()
print(f"Pure Python (list) time: {end - start:.4f} seconds")


start = time.perf_counter()
a_np + b_np
end = time.perf_counter()
print(f"NumPy (ndarray) time: {end - start:.4f} seconds")
```

```
#: python3 mat_add.py
Benchmarking matrix addition for 1000x1000 matrices...
Pure Python (list) time: 0.1569 seconds
NumPy (ndarray) time: 0.0026 seconds
```




Stack

- ❑ A collection of elements with Last-In-First-Out (LIFO) order.
- ❑ Push: Add an element to the top of the stack.
- ❑ Pop: Remove the top element from the stack.



Queue


- ❑ Similar as stack but with First-In-First-Out (FIFO) order.
- ❑ Enqueue: Add an element to the end of the queue.
- ❑ Dequeue: Remove the head element from the queue.



3. Trees and Graphs


Trees

- ❑ Represent a hierarchical relationship among data units.
- ❑ There is only one root node.
- ❑ Each node may have one or more children except the leaf nodes.



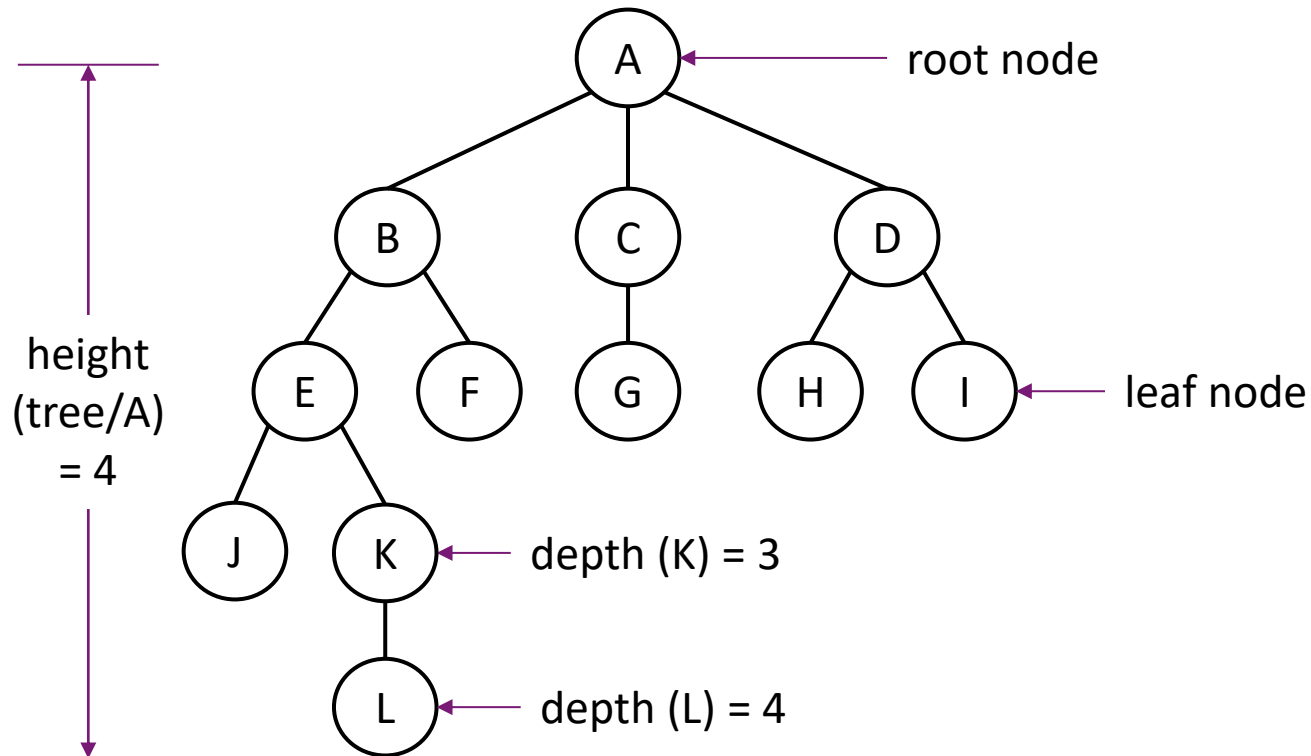
Relationships Among Nodes

- ❑ **Father node:** The direct parent of a node.
- ❑ **Sibling:** Nodes that share the same parent.
- ❑ **Child:** A node directly below a given node.
- ❑ **Ancestor:** A node above a given node.
- ❑ **Descendant:** A node below a given node.




Terminologies

- ❑ Height of a node: The longest path (no. of edges) from it to a leaf.
- ❑ Height of a tree: Height of the root node.
- ❑ Depth of a node: Number of edges from the root to the node.
- ❑ Degree of a node: Number of children of the node.




Binary Tree

- ❑ Binary tree: The degree of each node on a tree is at most two.
- ❑ Full binary tree: The degree of each node on a tree is two or zero.



Binary Tree



Full Binary Tree

Tree Traversal

❑ **Pre-order Traversal:** root => left subtree => right subtree


➤ A=>B=>E=>J=>K=>L=>F =>C=>G

❑ **Post-order Traversal:** left subtree => right subtree => root

➤ J=>L=>K=>E=>F=>B=>G=>C=>A

❑ **In-order Traversal:** left subtree => root => right subtree

➤ J=>E=>K=>L=>B=>F=>A=>C=>G



Implement Binary Tree with Python

❑ A binary tree has multiple nodes.

- One root node.
- Root and intermediate nodes have at most two children.
- Leaf nodes have no children.

❑ Using Optional[X] type: The value can be either of type X, or None.

```
@dataclass
class Node():
    value: int
    left: Optional[Node[int]] = None
    right: Optional[Node[int]] = None

@dataclass
class BinaryTree():
    root: Optional[Node[int]] = None
```

Implement Binary Tree with Python

❑ **Generic type:** Write one class or function that works for many types.

```
T = TypeVar("T")

@dataclass
class Node(Generic[T]):
    value: T
    left: Optional[Node[T]] = None
    right: Optional[Node[T]] = None

@dataclass
class BinaryTree(Generic[T]):
    root: Optional[Node[T]] = None
```


Question

☐ How to define the data structure of a tree?

➤ The number of child nodes are unlimited.

Implement Tree with Python

```
from __future__ import annotations
from dataclasses import dataclass, field

@dataclass
class TreeNode(Generic[T]):
    value: T
    # Way 1: Not allowed (assigning a mutable object: list).
    # children: List['TreeNode[T]'] = []
    # Way 2: List should be crated before use.
    # children: Optional[List['TreeNode[T]']] = None
    # Way 3: crate a list for each class instance via the factory.
    children: List[TreeNode[T]] = field(default_factory=list)

@dataclass
class Tree(Generic[T]):
    root: Optional[TreeNode[T]] = None
```

Practice (assisted with LLM)

❑ Implement a (binary) tree with the following features:


- Add: insert a node into the tree
- Remove: delete a node from the tree
- Traversal: visit all nodes using DFS or BFS
- Search: check for the existence of a node

Graphs


❑ Similar to trees except that:

- Graph may have loops, while trees cannot.
- There is no partial order (*e.g.*, parent vs child nodes) among nodes.

❑ A graph can be directed or undirected.



undirected graph




directed graph

Representing Undirected Graphs: Adjacent Matrix

□The adjacency matrix for undirected graph is defined as:

$$adj(i, j) = \begin{cases} 1, & \text{if there is an edge between vertex } i \text{ and } j \\ 0, & \text{otherwise} \end{cases}$$

□The adjacency matrix for undirected graphs is symmetric.



undirected graph


	A	B	C	D	E
A	0	0	1	1	0
B	0	0	1	0	0
C	1	1	0	1	1
D	1	0	1	0	1
E	0	0	1	1	0

adjacency matrix

Representing Directed Graphs: Adjacent Matrix

□The adjacency matrix for directed graph is defined as:

$$adj(i, j) = \begin{cases} 1, & \text{if there is an edge from } i \text{ to } j \\ 0, & \text{otherwise} \end{cases}$$



directed graph


	A	B	C	D	E
A	0	0	0	1	0
B	0	0	1	0	0
C	1	1	0	0	0
D	0	0	1	0	1
E	0	0	1	0	0

adjacency matrix

Weighted Graphs

□The adjacency matrix for undirected weighted graph is defined as:

$$adj(i, j) = \begin{cases} w, & \text{weight between node } i \text{ and } j \\ \infty, & \text{no edge between } i \text{ and } j \end{cases}$$



weighted graph

	A	B	C	D	E
A	∞	∞	2	3	∞
B	∞	∞	1	∞	∞
C	2	1	∞	5	2
D	3	∞	5	∞	4
E	∞	∞	2	4	∞


adjacency matrix

Graph Traversal: Two Orders


❑ **Depth-first Traversal:** Explore as far as possible along each branch before backtracking.

➤ A=>D=>C=>B=>C(visited) => back to B=>back to C=>back to D=>E=>C(visited)=>back to D=>back to A


❑ **Breadth-first Traversal:** Explore all the nodes at the present depth level before moving on to nodes at the next depth level.



directed graph



depth-first



breadth-first



Question

❑ How to define the data structure of a graph?

- 1) Unweighted vs Weighted graph.
- 2) Directed vs Undirected.

Implement Unweighted Graph via Adjacent Matrix

```
@dataclass
class Graph():
    adj: Dict[int, Set[int]] = field(default_factory=dict)
```

❑ Using data class with attributes (eq=True, frozen=True):

- eq: Two class instances can be compared by value.
 - This is required by dict.
- frozen=True: The dataclass is read only.

```
@dataclass(eq=True, frozen=True)
class Node():
    id: int

@dataclass
class Graph():
    adj: Dict[Node, Set[Node]] = field(default_factory=dict)
```

Implement Weighted Graph via Adjacent Matrix

```
@dataclass
class Node():
    id: int

@dataclass(eq=True, frozen=True)
class Graph():
    adj: Dict[Node, Dict[Node, int]] = field(default_factory=dict)
```

Implement Graph via List

```
@dataclass
class Node():
    id: int
    # unweighted version
    neighbors: List["Node"] = field(default_factory=list)
    # weighted version:
    List[Tuple["Node", int]] = field(default_factory=list)
```

```
@dataclass
class Graph():
    nodes: List[Node] = field(default_factory=list)
```

Another More Readable but Less Efficient Way

```
@dataclass
class Node():
    id: int

@dataclass
class Edge():
    u: Node
    v: Node
    weight: int

@dataclass
class Graph():
    nodes: List[Node] = field(default_factory=list)
    edges: List[Edge] = field(default_factory=list))
```

Practice (assisted with LLM)

❑ Implement an undirected/directed graph with the following features:

- Add: insert nodes and edges into the graph.
- Remove: delete nodes and edges from the graph.
- Traversal: visit all nodes using depth-first or breadth-first order.
- Search: check for the existence of nodes, edges, or paths.

4. Hash Set and Hash Map

Hash Set

- ❑ Python built-in type Set is a hash set.
- ❑ Map each value to an index using a hash function.
- ❑ Allow fast insertion, deletion, and lookup.
- ❑ The search time is constant.
- ❑ Tradeoff between space and time.

Index	Value
...	
231	JPM
...	
286	AAPL
...	
295	META
...	
...	

Index Calculation via Hash Function

□ A hash function takes some input and outputs a fixed-size integer.

Key
AAPL
MSFT
AMZN
META
TSLA
JPM



Index
286
314
310
295
308
231

Toy function: sum of ASCII Code

Example: AAPL

'A' = 65

'A' = 65

'P' = 80

'L' = 76

Sum: $65+65+80+76 = 286$

Collision and Handling

❑ **Collision:** When two different keys produce the same hash value.

➤ They are mapped to the same slot in the underlying array.

❑ **Possible handling strategy:**

➤ Separate chaining:

- Each slot in the table stores a linked list.
- When a collision occurs, simply append the new element to the list.

➤ Open addressing (Probing):

- Search for the next available slot using a probing sequence, *e.g.*, plus an offset.

Hash Table Provided in Python: Set

❑ For unique elements collection.

```
s = set()
s.add("apple")
s.add("banana")
s.add("orange")

s.remove("banana") # Remove elements
print("apple" in s) # True

for fruit in s:
    print(fruit)
```

Hash Map

- ❑ Python built-in type Dict is a hash map.
- ❑ Hash map is a data structure that stores key-value pairs.

Index	Key	Value
...		
231	JPM	
...		
286	AAPL	
...		
295	META	
...		
...		

Hash Map Provided in Python: Dict

❑ With both key and value.

```
students = {  
    "Alice": 20,  
    "Bob": 21,  
    "Charlie": 19  
}  
students["Dylan"] = 22  
del students["Alice"]  
print(students["Bob"]) # Output: 21
```

Practice (assisted with LLM)

□ Design and implement a Hash Map:

- Try different solutions for collision handling.
- Benchmark the performance.