

MF20006: Introduction to Computer Science

Lecture 7: Algorithm II

Hui Xu

xuh@fudan.edu.cn



Outline

1. Max/Min Heap

2. Binary Search Tree

3. Problems on Graphs



1. Max/Min Heap

Problem

- The stock exchange problem involves handling the concurrent arrival of ask (sell) and bid (buy) orders for various stocks.
- Ask order: The seller specifies the min price they are willing to sell.
- Bid order: The buyer specifies the max price they are willing to pay.

Bid/Ask	Capital	Change		
All-Hours ⓘ			...	☰
03:18 115.810	50 ▼	117.500	‑	1.27%
03:18 115.810	29 ▲	117.450	+	1.44%
03:19 115.800	25 ▲	117.440	+	0.83%
03:19 115.800	225 ▲	117.420	‑	0.94%
03:19 115.800	277 ▲	117.400	‑	3.22%
03:19 115.800	1 ▲	117.390	‑	2.39%
03:19 115.800	29 ▲	117.380	‑	1.81%
03:19 115.800	25 ▲	117.370	‑	1.26%
03:19 115.800	100 ▲	117.360	‑	1.11%
03:19 115.800	86 ▲	117.350	‑	2.30%
03:19 115.800	414 ▲	117.340	‑	1.40%
Bid		Ask		⾦
30.37%	██████████	69.63%		



Principle of Order Matching

- Orders are matched according to the price-time priority:
- Price priority: Higher bid prices have precedence over lower ones, and lower ask prices have precedence over higher ones.
- Time priority: If multiple orders have the same price, the one that was received earlier takes precedence.
- A trade is executed if the bid price is greater than or equal to the ask price.

Bid/Ask	Capital	Change		
All-Hours ⓘ			...	⋮
03:18 115.810	50 ▼	117.500	▬	1.27%
03:18 115.810	29 ▲	117.450	▬	1.44%
03:19 115.800	25 ▲	117.440	▬	0.83%
03:19 115.800	225 ▲	117.420	▬	0.94%
03:19 115.800	277 ▲	117.400	▬	3.22%
03:19 115.800	1 ▲	117.390	▬	2.39%
03:19 115.800	29 ▲	117.380	▬	1.81%
03:19 115.800	25 ▲	117.370	▬	1.26%
03:19 115.800	100 ▲	117.360	▬	1.11%
03:19 115.800	86 ▲	117.350	▬	2.30%
03:19 115.800	414 ▲	117.340	▬	1.40%
Bid		Ask	🏦	
30.37%	▬	69.63%	▬	

Trade Price

□ The trade price is typically determined based on the existing order.

□ At the ask price:

- There are ask orders in the book (with prices higher than existing bid orders).
- When a new bid order arrives with a price higher than or equal to the best ask, a trade is executed at the ask price.

□ At the bid price:

- There are bid orders in the book (with prices lower than existing ask orders).
- When a new ask order arrives with a price lower than or equal to the best bid, a trade is executed at the bid price.



Problem

□ Design an algorithm that efficiently manages incoming orders in real time, ensuring that trades are executed according to predefined matching rules.

□ Input:

- A continuous stream of ask and bid orders arriving concurrently.
- Each order specifies the quantity of shares and the price per share.

□ Output:

- Two dynamically maintained data structures, one for asks and one for bids.
- Each structure is sorted by price and arrival time for orders with the same price.
- These structures should support efficient order matching.




Sorting?

- Cost of retrieval the top- k elements is constant.
- However, the cost of inserting an element is $n/2$.




Max Heap

- A max heap is a type of binary tree.
- Heapify: Every node's value is greater than or equal to the values of its children.
- Completeness: Every level of the tree is filled from left to right.




Insert an Element

- To maintain completeness: Add a new value at the end.
- Heapify: Swap with parent while larger (“bubble up”).




Delete an Element

- Replace it with the last element (to maintain the complete tree).
- Remove the last element.
- Restore the heap property by either:
 - Heapify down towards the larger child if the new value is smaller than children.
 - Heapify up if the new value is larger than parent.







Find the Top-K Element

- Top 1: Just the root.** (20)
- Top 2: One of the children of the root.** (20) (16)
- Top 3: One of the children of top-2 elements.** (20) (16) (15)
- Top 4: ...**





Find the Top-K Element

- In the case of trading, we should delete the matched elements.
- Top k: Repeatedly remove the root k times



Implementation based on List

- Store the tree nodes in level-order traversal (top-down) as a list.
- The index of the children node given a node i:
 - Left: $2*i + 1$
 - Right: $2*i + 2$
- The index of the parent node given a node i:
 - $(i-1)/2$



level-order traversal

Proof?


□ Suppose the number of elements in the parent layer is n .

➤ The index should be $[n-1, \dots, 2n-2]$.

➤ Because the number of elements in previous layers is $n-1$.

□ The number of elements (if full) in the child layers is $2n$.

➤ The index should be $[2n-1, \dots, 4n-2]$



$n-1$ elements in previous layers

n elements in this layer

$2n$ elements in the child layer

Implementation

```
@dataclass
class MaxHeap:
    data: List[int] = field(default_factory=list)

    def push(self, val: int):
        """Insert a new value and bubble up."""
        self.data.append(val)
        i = len(self.data) - 1
        # Bubble up
        while i > 0:
            parent = (i - 1) // 2
            if self.data[i] > self.data[parent]:
                swap(self.data[i], self.data[parent])
                i = parent
            else:
                break
```



Implementation


```
def pop(self) -> int:  
    """Remove and return the max element."""  
    if not self.data:  
        raise IndexError("pop from empty heap")  
    n = len(self.data)  
    self.data[0], self.data[-1] = self.data[-1], self.data[0]  
    max_val = self.data.pop() # pop self.data[-1]  
    # Bubble down  
    i = 0  
    while True:  
        left, right = 2 * i + 1, 2 * i + 2  
        largest = i  
        if left < len(self.data) and self.data[left] > self.data[largest]:  
            largest = left  
        if right < len(self.data) and self.data[right] > self.data[largest]:  
            largest = right  
        if largest == i:  
            break  
        self.data[i], self.data[largest] = self.data[largest], self.data[i]  
        i = largest  
    return max_val
```

2. Binary Search Tree

Binary Search Tree

□ A binary search tree has some special rules to organize the data:


- Each node has two children at most: left and right.
- All nodes in the left subtree of a node are less than that node.
- All nodes in the right subtree of a node are greater than that node.



Data Structure of a BST


```
@dataclass
class Node:
    key: int
    left: Optional[Node] = None
    right: Optional[Node] = None

@dataclass
class BinarySearchTree:
    root: Optional[Node] = None
```



Construct a BST: Insert Nodes


```
def insert(self, key):
    new_node = Node(key)
    if self.root is None:
        self.root = new_node
        return
    current = self.root
    while True:
        if key < current.key:
            if current.left is None:
                current.left = new_node
                return
            current = current.left
        elif key > current.key:
            if current.right is None:
                current.right = new_node
                return
            current = current.right
        else:
            return
```



Construct a BST: Example


□ The order of arrival affects the tree structure.

Suppose the following arrival order:



Search on a BST



```
def search(self, key):
    current = self.root
    while current is not None:
        if key == current.key:
            return True
        elif key < current.key:
            current = current.left
        else:
            current = current.right
    return False
```



Cost of Operations on a BST


□ Search/Insertion/deletion: height of the tree

- Average: $O(\log n)$
- If the tree is highly imbalanced (worst case): $O(n)$




Self-Balancing BST: AVL Tree

- The heights of the two child subtrees of any node differ by at most one.
- Maintain a balancing factor for each node:
 - $BF(x) = \text{Height}(x \rightarrow \text{left}) - \text{Height}(x \rightarrow \text{right})$
- Rotate the subtree if $|BF(x)| > 1$




Imbalance Handling for Pattern 1: Rotation




The **left** subtree of imbalanced node **6** is higher.

The **left** subtree of node **2** is higher.




Imbalance Handling for Pattern 2: Rotation



The **left** subtree of imbalanced node **6** is higher.



The **right** subtree of node **2** is higher.



Still Imbalance



Imbalance Handling for Pattern 2: Rotation



Rotation for Different Imbalance Patterns

	Which subtree of the imbalanced node is higher?	Which child of the imbalanced node has a higher subtree?	Rotation Operation
Pattern 1	left	left	right rotation
Pattern 2	left	right	left-right rotation
Pattern 3	right	right	left rotation
Pattern 4	right	left	right-left rotation



Implementation with Python

```
@dataclass
class AVLNode:
    key: int
    left: Optional[AVLNode] = None
    right: Optional[AVLNode] = None
    height: int = 1

@dataclass
class AVLTree:
    root: Optional[AVLNode] = None
```



Implementation with Python


```
def insert(self, key):
    ... # insert node
    self._update_height(node)
    balance = self._balance_factor(node)
    # Left heavy
    if balance > 1:
        if key < node.left.key:
            return self._rotate_right(node) # LL
        else:
            node.left = self._rotate_left(node.left) # LR
            return self._rotate_right(node)
    # Right heavy
    if balance < -1:
        if key > node.right.key:
            return self._rotate_left(node) # RR
        else:
            node.right = self._rotate_right(node.right) # RL
            return self._rotate_left(node)
    return node
```



Red-black Tree


□ A BST with extra properties to ensure it stays roughly balanced.

- Every node is either red or black.
- The root is black.
- Every leaf (NIL / null pointer) is black.
- If a node is red, then both its children are black.
- For each node, all paths from the node to its descendant leaves have the same number of black nodes.



Red-black Tree: Insert a Node

□ Init as red.




Red-black Tree: Insert a Node

□ Local adjust based on some rules (rotations and recoloring).

➤ We do not enumerate all details.

□ Without needing to explicitly check the black-height of all paths.

➤ Compared to AVL: AVL tree needs to update the balance factor of ancestor nodes.




3. Problems on Graphs

Shortest Path Problem

□ Finding a path between two vertices in a graph such that the sum of the weights of its constituent edges is minimized.


➤ Unweighted graph: minimum number of edges.

➤ Weighted graph: accumulated weight of all edges.



unweighted graph

shortest path: A-D-E (cost: 2)




weighted graph

shortest path: A-B-D-E (cost: 6)




Dijkstra's Algorithm

- Each time select a node with the small distance from visited nodes.
- Maximize the number of nodes that can be visited given a budget.




	A	B	C	D	E
Visited?	Y	N	N	N	N
Distance	0	2	∞	5	∞




	A	B	C	D	E
Visited?	Y	Y	N	N	N
Distance	0	2	3	4	∞
Previous		A	B		

Dijkstra's Algorithm



	A	B	C	D	E
Visited?	Y	Y	Y	N	N
Distance	-	2	3	4	7
Previous		A	B	B	C



	A	B	C	D	E
Visited?	Y	Y	Y	Y	N
Distance	-	2	3	4	6
Previous		A	B	B	D

Complexity of Dijkstra's Algorithm

□ Supposing there are N nodes

- We outdated each node in each round
- In each round, we calculate the distance of the node to the rest nodes

□ Complexity: $O(n^2)$



What if the Graph Have Negative Weights?

□ Update the distance of visited nodes.

➤ Bellman–Ford algorithm




Minimal Spanning Tree


- A spanning tree of a connected graph is a subset of the edges that:
 - Connects all the vertices of the graph
 - Contains no cycles (so it's a tree)
- A minimum spanning tree is a spanning tree whose sum of edge weights is minimal among all possible spanning trees of the graph.



Example



Undirected Weighted Graph



Minimal Spanning Tree

Kruskal's Algorithm (greedy)

□ Steps:


- Sort all edges by weight (from smallest to largest).
- Start with an empty set of edges.
- Add edges one by one, skipping any that would form a cycle.
- Stop when you have $n-1$ edges.

□ Efficient for sparse graphs.

□ Uses Union-Find (Disjoint Set Union) to detect cycles.




Sort Edges by Weight



Undirected Weighted Graph

Weight	Edges
1	B-C
2	A-B, C-D
3	B-D
4	C-E
5	A-D, D-E
6	A-E

Add Edges One by One and Avoid Cycle



Weight	Edges
1	B-C
2	A-B, C-D
3	B-D
4	C-E
5	A-D, D-E
6	A-E

How to Detect Cycles?

□ Recognize connected components

- Union-find algorithm

□ Detect cycles for each connected component via DFS-based method:

- Start a Depth-First Search (DFS).
- Keep track of the parent node (the vertex you came from).
- If you encounter a neighbor that has already been visited and it is not the parent, then a cycle exists.



Union-Find Algorithm

- **Find:** Determine which set or connected component an element belongs to.
- **Union:** Merge two sets into a single set.

Edge	Connected Component
{A, B}	{A, B}
{A, C}	{A, B, C}
{B, C}	{A, B, C}
{B, H}	{A, B, C, H}
{D, E}	{A, B, C, H} {D, E}
{D, F}	{A, B, C, H} {D, E, F}
{E, F}	{A, B, C, H} {D, E, F}
{G, H}	{A, B, C, H, G} {D, E, F}



Prim's Algorithm (greedy)

□ Steps:


- Start from any vertex.
- Repeatedly add the smallest edge that connects a visited vertex to an unvisited vertex.
- Continue until all vertices are included.

□ Efficient for dense graphs.


□ Use a priority queue (min-heap) to pick edges quickly.



Example






Start from Vertex A



Add the smallest edge
from a visited vertex to an unvisited vertex

Example



Exercise: Vibe Coding

- Design experiments to benchmark the performance of naïve BST, AVL tree, and read-black tree.
- Design experiments to benchmark the performance of Kruskal's algorithm and Prim's algorithm.

